

HMM: Hidden Markov Models

Jean-Cédric Chappelier

Jean-Cedric.Chappelier@epfl.ch

Artificial Intelligence Laboratory

Objectives of this lecture

- ➔ Introduce **fundamental concepts** necessary to use HMMs for **PoS tagging**

Example: PoS tagging with HMM

Sentence to tag: **Time flies like an arrow.**

Example of HMM model:

□ PoS tags: $\mathcal{T} = \{\text{Adj}, \text{Adv}, \text{Det}, \text{N}, \text{V}, \dots\}$

□ Transition probabilities:

$$P(\text{N}|\text{Adj}) = 0.1, P(\text{V}|\text{N}) = 0.3, P(\text{Adv}|\text{N}) = 0.01, P(\text{Adv}|\text{V}) = 0.005, \\ P(\text{Det}|\text{Adv}) = 0.1, P(\text{Det}|\text{V}) = 0.3, P(\text{N}|\text{Det}) = 0.5$$

(plus all the others, such that stochastic constraints are fulfilled)

□ Initial probabilities: $P_I(\text{Adj}) = 0.01, P_I(\text{Adv}) = 0.001, P_I(\text{Det}) = 0.1, \\ P_I(\text{N}) = 0.2, P_I(\text{V}) = 0.003$ (+...)

□ Words: $\mathcal{L} = \{an, arrow, flies, like, time, \dots\}$

□ Emission probabilities: $P(time|\text{N}) = 0.1, P(time|\text{Adj}) = 0.01, P(time|\text{V}) = 0.005, \\ P(flies|\text{N}) = 0.1, P(flies|\text{V}) = 0.01, P(like|\text{Adv}) = 0.005, P(like|\text{V}) = 0.1, \\ P(an|\text{Det}) = 0.3, P(arrow|\text{N}) = 0.5$ (+...)

In this example, $12 = 3 \cdot 2 \cdot 2 \cdot 1 \cdot 1$ analyses are possible, for example:

$$P(\textit{time}/N \textit{ flies}/V \textit{ like}/Adv \textit{ an}/Det \textit{ arrow}/N) = 1.13 \cdot 10^{-11}$$

$$P(\textit{time}/Adj \textit{ flies}/N \textit{ like}/V \textit{ an}/Det \textit{ arrow}/N) = 6.75 \cdot 10^{-10}$$

$$\begin{aligned} &P(\textit{time}/N \textit{ flies}/V \textit{ like}/Adv \textit{ an}/Det \textit{ arrow}/N) \\ &= P_I(N) \cdot P(\textit{time}|N) \cdot P(V|N) \cdot P(\textit{flies}|V) \cdot P(Adv|V) \cdot P(\textit{like}|Adv) \\ &\quad \cdot P(Det|Adv) \cdot P(\textit{an}/Det) \cdot P(N|Det) \cdot P(\textit{arrow}|N) \\ &= 2e-1 \cdot 1e-1 \cdot 3e-1 \cdot 1e-2 \cdot 5e-3 \cdot 5e-3 \cdot 1e-1 \cdot 3e-1 \cdot 5e-1 \cdot 5e-1 \end{aligned}$$

The aim is to choose the most probable tagging

Contents

- HMM models, three basic problems
- Forward-Backward algorithms
- Viterbi algorithm
- Baum-Welch algorithm

Markov Models

Markov model: a discrete-time stochastic process \mathbf{T} on $\mathcal{T} = \{t^{(1)}, \dots, t^{(m)}\}$ satisfying the *Markov property* (limited conditioning)

$$P(T_i | T_1, \dots, T_{i-1}) = P(T_i | T_{i-k}, \dots, T_{i-1})$$

k : *order* of the Markov model

In practice $k = 1$ (bigrams) or 2 (trigrams) rarely 3 or 4 (\rightarrow learning difficulties)

From a theoretical point of view: every Markov model of order k can be represented as another Markov model of order 1 (choose $Y_i = (T_{i-k+1}, \dots, T_i)$)

Vocable:

$$P(T_1, \dots, T_i) = P(T_1) \cdot P(T_2 | T_1) \cdot \dots \cdot P(T_i | T_{i-1})$$

initial probabilities transition probabilities

Hidden Markov Models (HMM)

What is hidden?

☞ The model itself (i.e. the state sequence)

What do we see then?

☞ An *observation* w related to the state (but not the state itself)

Formally:

- a set of states $\mathcal{T} = \{t^{(1)}, \dots, t^{(m)}\}$ PoS tags
- a transition probabilities matrix \mathbf{A} such that $A_{tt'} = P(T_{i+1} = t' | T_i = t)$, shorten $P(t'|t)$ (independant of i)
- an initial probabilities vector \mathbf{I} such that $I_t = P(T_1 = t)$, shorten $P_I(t)$
- ☆ an alphabet \mathcal{L} (not necessarily finite) words
- ☆ n probability densities on \mathcal{L} (*emission probabilities*): $B_t(w) = P(W_i = w | T_i = t)$ (for $w \in \mathcal{L}$), shorten $P(w|t)$.

Simple example of HMM

Example: a cheater tossing from two hidden (unfair) coins


States: coin 1 and coin 2: $\mathcal{T} = \{1, 2\}$

transition matrix $\mathbf{A} = \begin{bmatrix} 0.4 & 0.6 \\ 0.9 & 0.1 \end{bmatrix}$

$$\text{alphabet} = \{H, T\}$$

emission probabilities $\mathbf{B}_1 = (0.49, 0.51)$ et $\mathbf{B}_2 = (0.85, 0.15)$

initial probabilities $\mathbf{I} = (0.5, 0.5)$



👉 5 free parameters: $I_1, A_{11}, A_{21}, B_1(\text{H}), B_2(\text{H})$

Observation: HTTHTTHHTTHTTTTHHTTHTTTTHTHHTHTHHTTTT

(Hidden) sequence of states: 211211211121112111211121121121211112]

The three basic problems for HMMs

Problems: Given an HMM and an observation sequence $\mathbf{w} = w_1 \dots w_n$

- ⇒ given the parameters θ of the HMM, what is the probability of the observation sequence:

$$P(\mathbf{w}|\theta)$$

Application: Language Identification

- ⇒ given the parameters θ of the HMM, find the most likely state sequence \mathbf{t} that produces \mathbf{w} :

$$\underset{\mathbf{t}}{\operatorname{Argmax}} P(\mathbf{t}|\mathbf{w}, \theta)$$

Application: PoS Tagging, Speech recognition

- ⇒ find the parameters that maximize the probability of producing \mathbf{w} : $\underset{\theta}{\operatorname{Argmax}} P(\theta|\mathbf{w})$

Application: Unsupervised learning

Remarks:

- ❶ $\theta = (\mathbf{I}, \mathbf{A}, \mathbf{B})$
 $= (I_1, \dots, I_m, B_1(w_1), B_1(w_2), \dots, B_1(w_L), B_2(w_1), \dots, B_2(w_L),$
 $\dots, B_m(w_1), \dots, B_m(w_L), A_{11}, \dots, A_{1m}, \dots, A_{m1}, \dots, A_{mm})$
i.e. $(m - 1) + m \cdot (L - 1) + m \cdot (m - 1) = m \cdot (m + L - 1) - 1$ free
parameters (because of sum-to-1 constraints), where $m = |\mathcal{T}|$ and $L = |\mathcal{L}|$ (in the
finite case, otherwise L stands for the total number of parameters used to represent \mathcal{L})
- ❷ Supervised learning (i.e. $\underset{\theta}{\text{Argmax}} P(\theta|\mathbf{w}, \mathbf{t})$) is easy
- ❸ **WARNING!** There is a difference between $P(\theta|\mathbf{w})$ and $P(\mathcal{M}|\mathbf{w})$!
The model \mathcal{M} is supposed to be known here, but its parameters θ : i.e. the HMM
design is already defined (number of states, alphabet) only the parameters are missing.

Contents

- ➡ HMM models, three basic problems
- 👉 Forward-Backward algorithms
- ➡ Viterbi algorithm
- ➡ Baum-Welch algorithm

Computation of $P(\mathbf{W}|\boldsymbol{\theta})$

Computation of $P(\mathbf{W}|\boldsymbol{\theta})$ is mathematically trivial:

$$P(\mathbf{W}|\boldsymbol{\theta}) = \sum_{\mathbf{t}} P(\mathbf{W}, \mathbf{t}|\boldsymbol{\theta}) = \sum_{\mathbf{t}} P(\mathbf{W}|\mathbf{t}, \boldsymbol{\theta}) \cdot P(\mathbf{t}|\boldsymbol{\theta})$$

Practical limitation: complexity is $\mathcal{O}(n m^n)$ \rightsquigarrow exponential!

Practical computation: forward/backward algorithms \longrightarrow complexity is $\mathcal{O}(n m^2)$

Forward-Backward algorithms

"forward" variable : $\alpha_i(t) = P(w_1, \dots, w_i, T_i = t | \theta)$ $t \in \mathcal{T}$

iterative computation: $\alpha_{i+1}(t') = B_{t'}(w_{i+1}) \cdot \sum_{t \in \mathcal{T}} (\alpha_i(t) \cdot A_{tt'})$

$$\alpha_1(t) = B_t(w_1) \cdot I_t$$

"backward" variable : $\beta_i(t) = P(w_{i+1}, \dots, w_n | T_i = t, \theta)$

$$\beta_{i-1}(t') = \sum_{t \in \mathcal{T}} (\beta_i(t) \cdot A_{t't} \cdot B_t(w_i))$$

$$\beta_n(t) = 1 \text{ (by convention, practical considerations)}$$

Forward-Backward algorithms (2)

"forward-backward" variable : $\gamma_i(t) = P(T_i = t | \mathbf{w}, \boldsymbol{\theta})$

$$\gamma_i(t) = \frac{P(\mathbf{w}, T_i = t | \boldsymbol{\theta})}{P(\mathbf{w} | \boldsymbol{\theta})} = \frac{\alpha_i(t) \cdot \beta_i(t)}{\sum_{t' \in \mathcal{T}} \alpha_i(t') \cdot \beta_i(t')}$$

Computation in $\mathcal{O}(n m^2)$ \rightarrow efficient solutions to "first problem":

$$P(\mathbf{w} | \boldsymbol{\theta}) = \sum_{t \in \mathcal{T}} P(\mathbf{w}, T_n = t | \boldsymbol{\theta}) = \sum_{t \in \mathcal{T}} \alpha_n(t)$$

$$P(\mathbf{w} | \boldsymbol{\theta}) = \sum_{t \in \mathcal{T}} \alpha_i(t) \cdot \beta_i(t) \quad \forall i : 1 \leq i \leq n$$

Contents


- ➡ HMM models, three basic problems
- ➡ Forward-Backward algorithms
- 👉 Viterbi algorithm
- ➡ Baum-Welch algorithm

Viterbi algorithm (1)

Efficient solution to the "second problem": find the most likely sequence of states \mathbf{t} (knowing \mathbf{w} and the parameters $\boldsymbol{\theta}$) : $\text{Argmax}_{\mathbf{t}} P(\mathbf{t}|\mathbf{w}, \boldsymbol{\theta})$

\Rightarrow maximize (in \mathbf{t}) $P(\mathbf{t}, \mathbf{w}|\boldsymbol{\theta})$.

"The" lattice  temporal unfolding of all possible walks through the Markov chain



Viterbi algorithm (2)

Let $\rho_i(t) = \max_{t_1, \dots, t_{i-1}} P(t_1, \dots, t_{i-1}, t_i = t, w_1, \dots, w_i | \theta)$

We are looking for $\max_{t \in \mathcal{T}} \rho_n(t)$

It's easy (exercise) to show that $\rho_i(t) = \max_{t'} \left[P(t|t', \theta) P(w_i|t, \theta) \rho_{i-1}(t') \right]$

from which the following algorithm comes:

...

Viterbi algorithm (3)

for all $t \in \mathcal{T}$ do

$$\rho_1(t) = I_t \cdot B_t(w_1)$$

.....

for i from 2 to n do

for all $t \in \mathcal{T}$ do

- $\rho_i(t) = B_t(w_i) \cdot \max_{t'} (A_{t't} \cdot \rho_{i-1}(t'))$
- mark one of the transitions from t' to t where the maximum is reached

.....

reconstruct backwards (from t_n) the best path following the marked transitions

Contents

- ➡ HMM models, three basic problems
- ➡ Forward-Backward algorithms
- ➡ Viterbi algorithm
- 👉 Baum-Welch algorithm

Expectation-Maximization

Our goal: maximize $P(\boldsymbol{\theta}|\mathbf{w})$

☞ Maximum-likelihood estimation of $\boldsymbol{\theta}$ \rightarrow maximization of $P(\mathbf{w}|\boldsymbol{\theta})$

To achieve it: **Expectation-Maximization (EM) algorithm**

General formulation of EM:

given

- observed data $\mathbf{w} = w_1 \dots w_n$
- a parameterized probability distribution $P(\mathbf{T}, \mathbf{W}|\boldsymbol{\theta})$ where
 - $\mathbf{T} = T_1 \dots T_n$ are unobserved data
 - $\boldsymbol{\theta}$ are the parameters of the model

determine $\boldsymbol{\theta}$ that maximizes $P(\mathbf{w}|\boldsymbol{\theta})$ by convergence of iterative computation of the series $\boldsymbol{\theta}^{(i)}$ that maximizes (in $\boldsymbol{\theta}$)

$$\mathbf{E}_{\mathbf{T}} [\log P(\mathbf{T}, \mathbf{W}|\boldsymbol{\theta})|\mathbf{w}, \boldsymbol{\theta}^{(i-1)}]$$

Expectation-Maximization (2)

To do so, define the auxiliary function

$$Q(\boldsymbol{\theta}, \boldsymbol{\theta}') = \mathbf{E}_{\mathbf{T}} [\log P(\mathbf{T}, \mathbf{W} | \boldsymbol{\theta}) | \mathbf{w}, \boldsymbol{\theta}'] = \sum_{\mathbf{t}} P(\mathbf{t} | \mathbf{w}, \boldsymbol{\theta}') \log P(\mathbf{t}, \mathbf{w} | \boldsymbol{\theta})$$

as it can be shown (with Jensen inequality) that

$$Q(\boldsymbol{\theta}, \boldsymbol{\theta}') > Q(\boldsymbol{\theta}', \boldsymbol{\theta}') \Rightarrow P(\mathbf{w} | \boldsymbol{\theta}) > P(\mathbf{w} | \boldsymbol{\theta}')$$

This is the fundamental principle of EM: **if** we already have an estimation $\boldsymbol{\theta}'$ of the parameters and we find another parameter configuration $\boldsymbol{\theta}$ for which the first inequality (on Q) holds, **then** \mathbf{w} is most probable with model $\boldsymbol{\theta}$ rather than with model $\boldsymbol{\theta}'$.

Expectation-Maximization (3)

EM algorithm:

- Estimation Step: Compute $Q(\boldsymbol{\theta}, \boldsymbol{\theta}^{(i)})$
- Maximization Step: Compute $\boldsymbol{\theta}^{(i+1)} = \underset{\boldsymbol{\theta}}{\text{Argmax}} Q(\boldsymbol{\theta}, \boldsymbol{\theta}^{(i)})$

in other words:

1. Choose $\boldsymbol{\theta}^{(0)}$ (and set $i = 0$)
2. Find $\boldsymbol{\theta}^{(i+1)}$ which maximizes $\sum_{\mathbf{t}} P(\mathbf{t}|\mathbf{w}, \boldsymbol{\theta}^{(i)}) \log P(\mathbf{t}, \mathbf{w}|\boldsymbol{\theta}^{(i+1)})$
3. Set $i \leftarrow i + 1$ and go back to (2) unless some convergence test is fulfilled

Baum-Welch Algorithm

The Baum-Welch Algorithm is an EM algorithm for estimating HMM parameters. It's an answer to the "third problem".

The goal is therefore to find

$$\underset{\theta}{\operatorname{Argmax}} \sum_{\mathbf{t}} P(\mathbf{t}|\mathbf{w}, \theta') \log P(\mathbf{t}, \mathbf{w}|\theta) = \underset{\theta}{\operatorname{Argmax}} \underbrace{\sum_{\mathbf{t}} P(\mathbf{t}, \mathbf{w}|\theta') \log P(\mathbf{t}, \mathbf{w}|\theta)}_{\stackrel{\text{def}}{=} \hat{Q}(\theta, \theta')}$$

since $P(\mathbf{w}|\theta')$ does not depend on θ .

What is $\log P(\mathbf{t}, \mathbf{w}|\theta)$?

$$\log P(\mathbf{t}, \mathbf{w}|\theta) = \log P_I(t_1) + \sum_{i=2}^n \log P(t_i|t_{i-1}) + \sum_{i=1}^n \log P(w_i|t_i)$$

$\hat{Q}(\boldsymbol{\theta}, \boldsymbol{\theta}')$ consists therefore of 3 terms:

$$\hat{Q}((\mathbf{I}, \mathbf{A}, \mathbf{B}), \boldsymbol{\theta}') = Q_I(\mathbf{I}, \boldsymbol{\theta}') + Q_A(\mathbf{A}, \boldsymbol{\theta}') + Q_B(\mathbf{B}, \boldsymbol{\theta}')$$

Let's compute one of these:

$$\begin{aligned} Q_I(\mathbf{I}, \boldsymbol{\theta}') &= \sum_{\mathbf{t}} P(\mathbf{t}, \mathbf{w} | \boldsymbol{\theta}') \log P_I(t_1) \\ &= \sum_{t_1} \sum_{t_2, \dots, t_n} P(t_1, \mathbf{w} | \boldsymbol{\theta}') \cdot P(t_2, \dots, t_n | t_1, \mathbf{w}, \boldsymbol{\theta}') \cdot \log P_I(t_1) \\ &= \sum_{t \in \mathcal{T}} P(t_1 = t, \mathbf{w} | \boldsymbol{\theta}') \cdot \log P_I(t) \underbrace{\sum_{t_2, \dots, t_n} P(t_2, \dots, t_n | t_1, \mathbf{w}, \boldsymbol{\theta}')}_{=1} \\ &= \sum_{t \in \mathcal{T}} P(t_1 = t, \mathbf{w} | \boldsymbol{\theta}') \cdot \log I_t \end{aligned}$$

Similarly we have:

$$Q_A(\mathbf{A}, \boldsymbol{\theta}') = \sum_{i=2}^n \sum_{t \in \mathcal{T}} \sum_{t' \in \mathcal{T}} P(T_{i-1} = t, T_i = t', \mathbf{w} | \boldsymbol{\theta}') \log A_{tt'}$$

$$Q_B(\mathbf{B}, \boldsymbol{\theta}') = \sum_{i=2}^n \sum_{t \in \mathcal{T}} P(T_i = t, \mathbf{w} | \boldsymbol{\theta}') \log B_t(w_i)$$

Therefore \hat{Q} is a sum of three **independent** terms (e.g. Q_I does not depend on \mathbf{A} nor on \mathbf{B}) and therefore the maximisation over $\boldsymbol{\theta}$ is achieved by **the three terms separately**, i.e. maximizing $Q_I(\mathbf{I}, \boldsymbol{\theta}')$ over \mathbf{I} , $Q_A(\mathbf{A}, \boldsymbol{\theta}')$ over \mathbf{A} and $Q_B(\mathbf{B}, \boldsymbol{\theta}')$ over \mathbf{B} separately.

Notice that all these three functions are sums (over i) of functions of the form:

$$f(\mathbf{x}) = \sum_{j=1}^m y_j \log x_j$$

and all the above three functions have to be maximized under the constraint $\sum_{j=1}^m x_j = 1$.^a

This maximization under constraints is achieved using Lagrange multipliers, i.e. looking at

$$g(\mathbf{x}) = f(\mathbf{x}) - \lambda \cdot \sum_{j=1}^m x_j = \sum_{j=1}^m (y_j \log x_j - \lambda \cdot x_j)$$

Solving this by $\frac{\partial}{\partial x} g(x) = 0$, we find that $\lambda = \frac{y_j}{x_j}$.

Putting this back in the constraint we find:

$$x_j = \frac{y_j}{\sum_{j=1}^m y_j}$$

^aTo be accurate: for \mathbf{B}_t the constraint is $\sum_{w \in \mathcal{L}} B_t(w) = 1$. This changes the formulas a bit, but not the essence of the computation.

Summarizing the obtained results, we have the following reestimation formulas (where the max. is reached):

$$\hat{I}_t = \frac{P(T_1 = t, \mathbf{w} | \boldsymbol{\theta}')}{\sum_{t' \in \mathcal{T}} P(T_1 = t', \mathbf{w} | \boldsymbol{\theta}')} = \frac{P(T_1 = t, \mathbf{w} | \boldsymbol{\theta}')}{P(\mathbf{w} | \boldsymbol{\theta}')}$$

$$\begin{aligned} \widehat{A_{tt'}} &= \frac{\sum_{i=2}^n P(T_{i-1} = t, T_i = t', \mathbf{w} | \boldsymbol{\theta}')}{\sum_{i=2}^n \sum_{\tau \in \mathcal{T}} P(T_{i-1} = t, T_i = \tau, \mathbf{w} | \boldsymbol{\theta}')} \\ &= \frac{\sum_{i=2}^n P(T_{i-1} = t, T_i = t', \mathbf{w} | \boldsymbol{\theta}')}{\sum_{i=2}^n P(T_{i-1} = t, \mathbf{w} | \boldsymbol{\theta}')} \end{aligned}$$

and:

$$\widehat{B_t(w)} = \frac{\sum_{\substack{i=1 \text{ s.t.} \\ w_i=w}}^n P(T_i = t, \mathbf{w}|\boldsymbol{\theta}')}{\sum_{i=2}^n P(T_i = t, \mathbf{w}|\boldsymbol{\theta}')} = \frac{\sum_{i=2}^n P(T_i = t, \mathbf{w}|\boldsymbol{\theta}') \delta_{w_i, w}}{\sum_{i=2}^n P(T_i = t, \mathbf{w}|\boldsymbol{\theta}')}$$

with $\delta_{w, w'} = 1$ if $w = w'$ and 0 otherwise.

Baum-Welch Algorithm: effective computation

How do we compute these reestimation formulas?

Let $\chi_i(t, t') = P(T_i = t, T_{i+1} = t' | \mathbf{w})$

χ_i is easy to compute with "forward" and "backward" variables:

$$\chi_i(t, t') = \frac{\alpha_i(t) \cdot A_{tt'} \cdot B_{t'}(w_{i+1}) \cdot \beta_{i+1}(t')}{\sum_{\tau \in \mathcal{T}} \sum_{\tau' \in \mathcal{T}} \alpha_i(\tau) \cdot A_{\tau\tau'} \cdot B_{\tau'}(w_{i+1}) \cdot \beta_{i+1}(\tau')}$$

Notice: $\gamma_i(t) = \sum_{t' \in \mathcal{T}} \chi_i(t, t')$ for all $1 \leq i < n$

Effective Reestimation formulas

$$\hat{I}_t = \gamma_1(t)$$

$$\widehat{A_{tt'}} = \frac{\sum_{i=1}^{n-1} \chi_i(t, t')}{\sum_{i=1}^{n-1} \gamma_i(t)}$$

$$\widehat{B_t(w)} = \frac{\sum_{\substack{i=1 \text{ s.t.} \\ w_i=w}}^n \gamma_i(t)}{\sum_{i=1}^n \gamma_i(t)} = \frac{\sum_{i=1}^n \gamma_i(t) \delta_{w_i, w}}{\sum_{i=1}^n \gamma_i(t)}$$

with $\delta_{w, w'} = 1$ if $w = w'$ and 0 otherwise.

Baum-Welch Algorithm

1. Let $\theta^{(0)}$ be an initial parameter set
2. Compute iteratively α , β and then γ and χ
3. Compute $\theta^{(t+1)}$ with reestimation formulas
4. If $\theta^{(t+1)} \neq \theta^{(t)}$, go to (2) [or another weaker stop test]

WARNING!

The algorithm converges but only towards a local maximum of $\mathbf{E} [\log P(\mathbf{T}, \mathbf{W}|\theta)]$


CRF versus HMM

(linear) **Conditional Random Fields** (CRF) are a discriminative generalization of the HMMs where “features” no longer needs to be state-conditionnal probabilities (less constraint).

For instance (order 1):

HMM


$$P(\mathbf{t}, \mathbf{w}) = P(t_1) P(w_1|t_1) \cdot \prod_{i=2}^n P(w_i|t_i) P(t_i|t_{i-1})$$



CRF

$$P(\mathbf{t}|\mathbf{w}) = \prod_{i=2}^n P(t_{i-1}, t_i|\mathbf{w})$$

(with $P(t_{i-1}, t_i|\mathbf{w}) \propto \exp \left(\sum_j \lambda_j f_j(t_{i-1}, t_i, \mathbf{w}, i) \right)$)



Keypoints

- ⇒ HMMs definitions, their applications
- ⇒ Three basic problems for HMMs
- ⇒ Algorithms needed to solve these problems: Forward-Backward, Viterbi, Baum-Welch (be aware of their existence, but not the implementation details)

References

- [1] L. R. Rabiner, *A Tutorial on Hidden Markov Models and Selected Applications in Speech Recognition*, Proceedings of the IEEE, vol. 77, No. 2, 1989.
- [2] C. D. Manning, H. Schütze, *Foundations of Statistical Natural Language Processing*, MIT, 1999.
- [3] A. P. Dempster, N. M. Laird, D. B. Rubin, *Maximum-likelihood from incomplete data via the EM algorithm*, Journal of Royal Statistical Society B, 1977.
- [4] H. Bourlard et al., *Traitement de la parole*, 2000; pp. 179-200, 202-214, 232-260.

APPENDUM

Justification of the maximization of the auxiliary function Q for finding θ maximizing $P(\mathbf{w}|\theta)$:

$$\begin{aligned}\log P(\mathbf{w}|\theta) - \log P(\mathbf{w}|\theta') &= \log \frac{P(\mathbf{w}|\theta)}{P(\mathbf{w}|\theta')} = \log \sum_{\mathbf{t}} \frac{P(\mathbf{w}, \mathbf{t}|\theta)}{P(\mathbf{w}|\theta')} \\&= \log \sum_{\mathbf{t}} P(\mathbf{t}|\mathbf{w}, \theta') \frac{P(\mathbf{w}, \mathbf{t}|\theta)}{P(\mathbf{w}, \mathbf{t}|\theta')} \\&\stackrel{\text{Jensen}}{\geq} \sum_{\mathbf{t}} P(\mathbf{t}|\mathbf{w}, \theta') \log \frac{P(\mathbf{w}, \mathbf{t}|\theta)}{P(\mathbf{w}, \mathbf{t}|\theta')} \\&\geq \mathbf{E}_{\mathbf{T}} [\log P(\mathbf{T}, \mathbf{W}|\theta) | \mathbf{w}, \theta'] - \mathbf{E}_{\mathbf{T}} [\log P(\mathbf{T}, \mathbf{W}|\theta') | \mathbf{w}, \theta'] \\&\geq Q(\theta, \theta') - Q(\theta', \theta')\end{aligned}$$

Therefore:

$$Q(\theta, \theta') > Q(\theta', \theta') \Rightarrow \log P(\mathbf{w}|\theta) > \log P(\mathbf{w}|\theta') \Rightarrow P(\mathbf{w}|\theta) > P(\mathbf{w}|\theta')$$