


# DEEP LEARNING CRASH COURSE

- Single Layer Perceptron
- Multiple Layer Perceptron
- Convolutional Neural Net



# ARTIFICIAL INTELLIGENCE




1997: Deep Blue beats chess World Champion




2016: AlphaGo beats go world champion

# LINEAR CLASSIFICATION




$$f(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{w} \cdot \mathbf{x} + \mathbf{b} \geq 0, \\ 0 & \text{otherwise.} \end{cases}$$

# SINGLE LAYER PERCEPTRON




$\sigma$ :

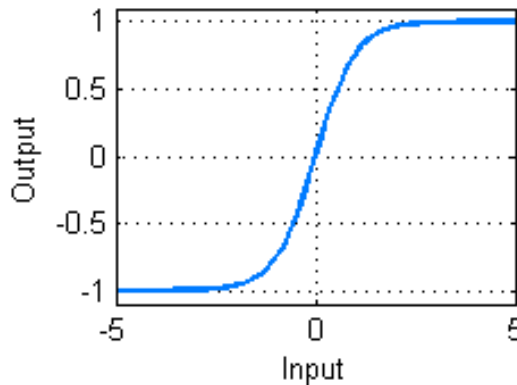
Threshold




Logistic Sigmoid




tanh



Hinge Loss



# MULTILAYER PERCEPTRON




$$\mathbf{h} = \sigma_1(\mathbf{W}_1\mathbf{x} + \mathbf{b}_1)$$

$$\mathbf{y} = \sigma_2(\mathbf{W}_2\mathbf{h} + \mathbf{b}_2)$$

- The process can be repeated several times to create a vector  $\mathbf{h}$ .



# MULTILAYER PERCEPTRON



$$\mathbf{h} = \sigma_1(\mathbf{W}_1\mathbf{x} + \mathbf{b}_1)$$

$$\mathbf{y} = \sigma_2(\mathbf{W}_2\mathbf{h} + \mathbf{b}_2)$$

- The process can be repeated several times to create a vector  $\mathbf{h}$ .
- It can then be done again to produce an output  $\mathbf{y}$ .
- —> This output is a **differentiable** function of the weights.

# BINARY CASE



Given a training set  $\{\mathbf{x}_n, t_n\}_{1 \leq n \leq N}$  where  $t_n \in \{0, 1\}$ , minimize

$$E(\mathbf{W}, \mathbf{b}) = -\frac{1}{N} \sum_1^N [t_i \log(y_i) + (1 - t_i) \log(1 - y_i)] ,$$

$$\begin{aligned} y_i &= f(\mathbf{x}_i) \\ &= \sigma(\mathbf{W}_2(\sigma(\mathbf{W}_1 \mathbf{x}_i + \mathbf{b}_1)) + \mathbf{b}_2), \end{aligned}$$

that is, minimize the number of misclassified samples.

Since  $E$  is a differentiable function of  $\mathbf{W}$  and  $\mathbf{b}$ , this can be done using a gradient-based technique, also known as back propagation.

# MULTI-CLASS CASE



In the multi-class case, the probability that input vector  $\mathbf{x}$  belongs to class  $c$  is taken to be

$$P(\mathbf{x} \in c | \mathbf{W}, \mathbf{b}) = \frac{y^c(\mathbf{x}; \mathbf{W}, \mathbf{b})}{\sum_k y^k(\mathbf{x}; \mathbf{W}, \mathbf{b})} .$$


Given a training set  $\{\mathbf{x}_n, t_n^1, \dots, t_n^C\}_{1 \leq n \leq N}$  where  $t_n^c \in \{0, 1\}$ , minimize

$$E(\mathbf{W}, \mathbf{b}) = - \sum_n \sum_c t_n^c \log(P(\mathbf{x}_n \in c | \mathbf{W}, \mathbf{b})) .$$

Since  $E$  remains a differentiable function of  $\mathbf{W}$  and  $\mathbf{b}$ , this can also be done using a gradient-based technique, also known as back propagation.



# HINGE LOSS OR RELU




$$\mathbf{h} = \sigma(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$$

$$\mathbf{y} = \sigma(\mathbf{W}_2 \mathbf{h} + \mathbf{b}_2)$$

$$\sigma(\mathbf{x}) = \max(0, \mathbf{x}).$$

- Each node defines a hyperplane.
- The resulting function is piecewise linear affine and continuous.


# ONE SINGLE HYPERPLANE



# TWO HYPERPLANES

$$\begin{cases} \mathbf{h} = \max(\mathbf{W}\mathbf{x} + \mathbf{b}, 0) \\ y = \mathbf{w}'^\top \mathbf{h} \end{cases}$$


with  $\dim(\mathbf{h}) = 2$




# THREE HYPERPLANES

$$\begin{cases} \mathbf{h} = \max(\mathbf{W}\mathbf{x} + \mathbf{b}, 0) \\ y = \mathbf{w}'^\top \mathbf{h} \end{cases}$$

with  $\dim(\mathbf{h}) = 3$





# SIGMOID AND TANH



$$\mathbf{h} = \sigma(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$$


$$\mathbf{y} = \sigma(\mathbf{W}_2 \mathbf{h} + \mathbf{b}_2)$$




$$\begin{aligned} \text{sigm: } \sigma(x) &= \frac{1}{1 + \exp(-x)} \\ \text{tanh: } \sigma(x) &= \frac{\exp(x) - \exp(-x)}{\exp(x) + \exp(-x)} \end{aligned}$$

- Each node defines a hyperplane.
- The resulting function is continuously differentiable.

# INTERPOLATING A SURFACE



$$z = 100 * (y - x^2)^2 + (1 - x)^2$$

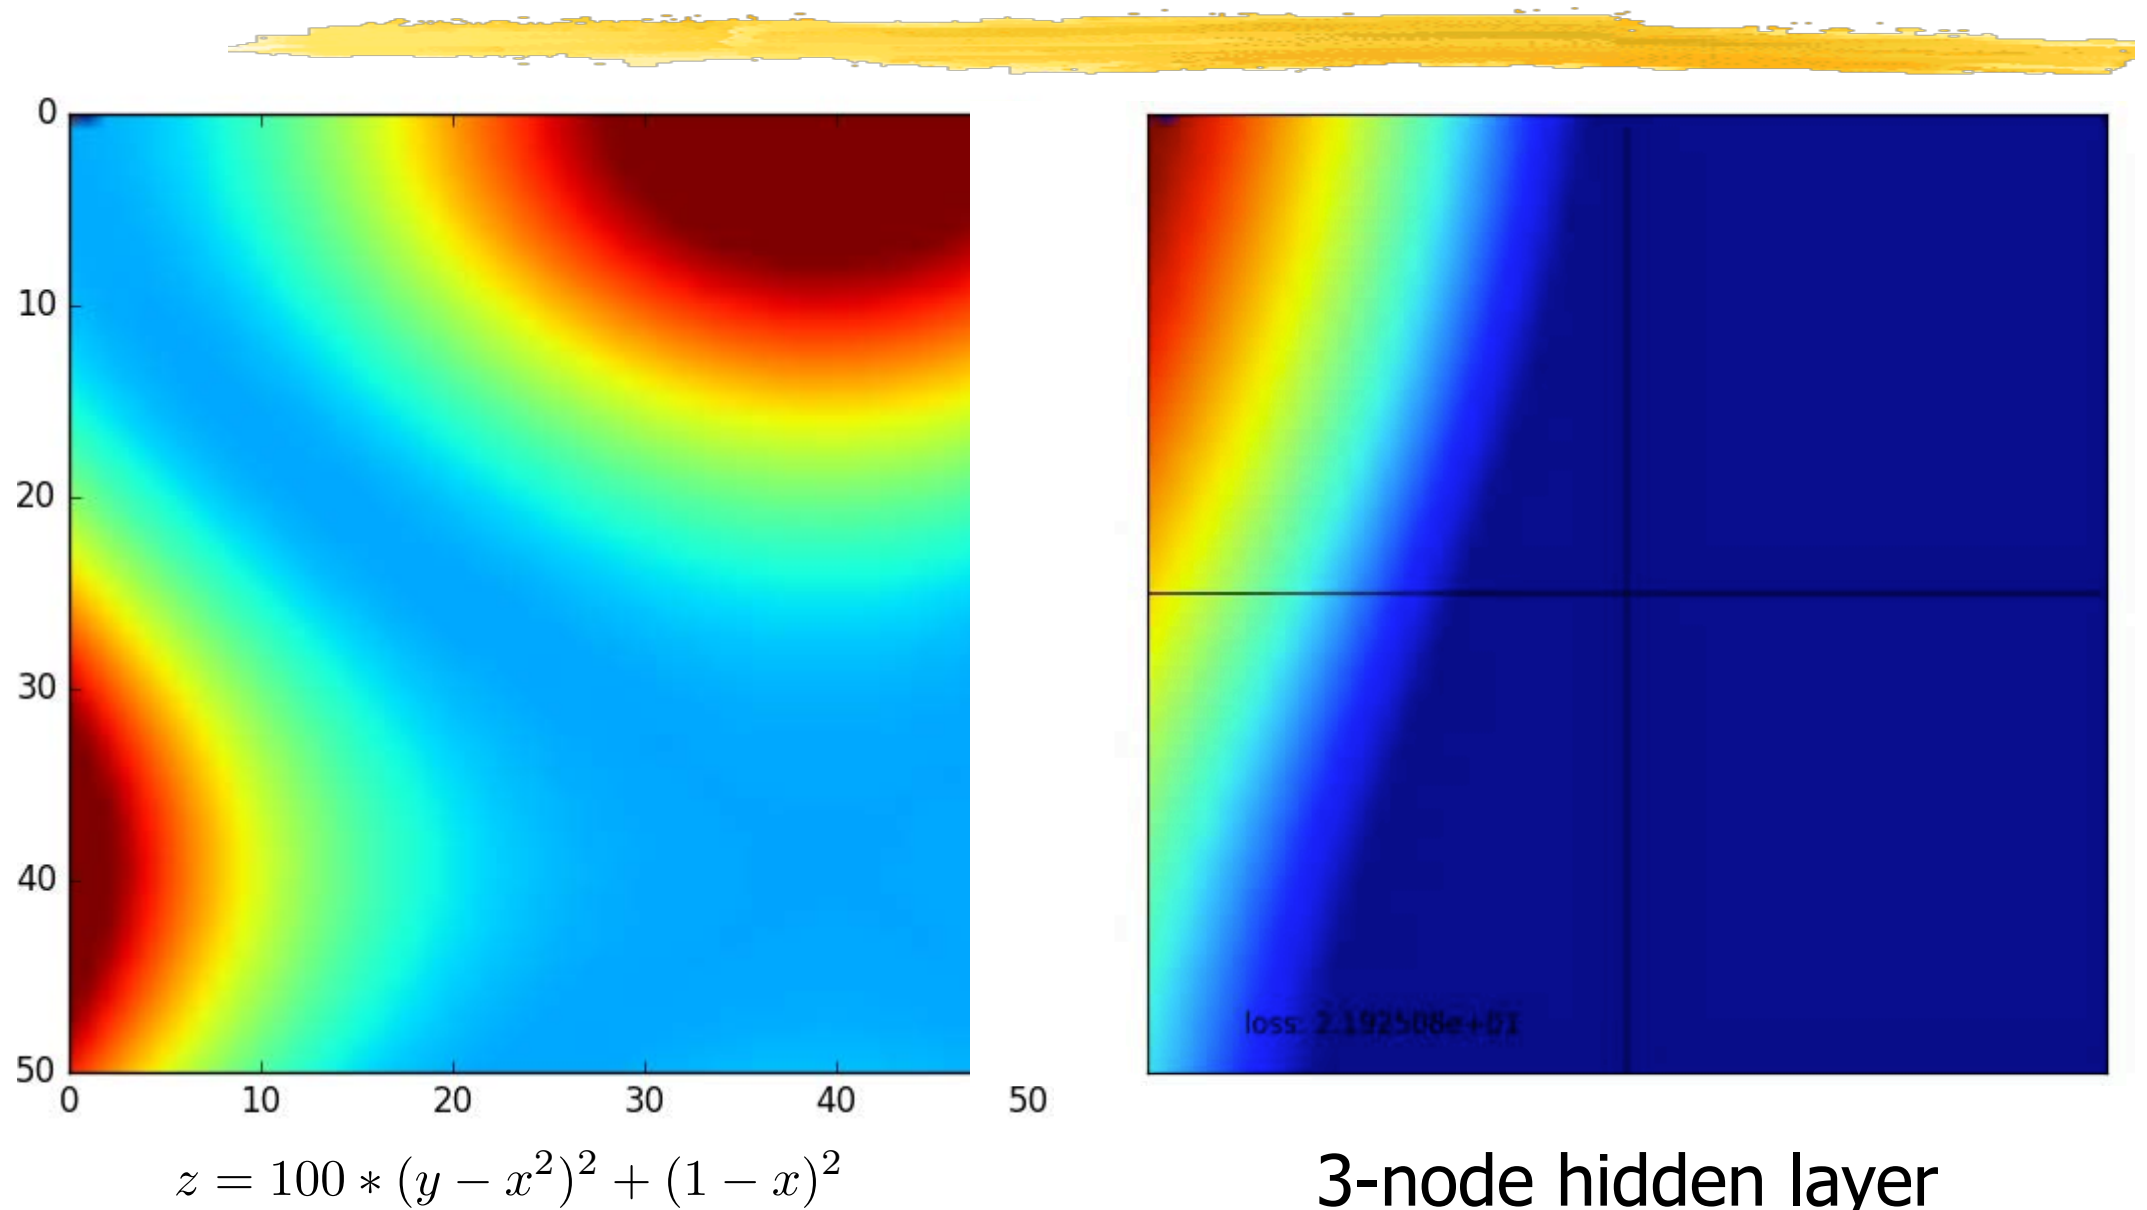


$$\begin{aligned} z &= f(x, y) \\ &= \mathbf{w}_2 \sigma \left( \mathbf{W}_1 \begin{bmatrix} x \\ y \end{bmatrix} + \begin{bmatrix} b_x \\ b_y \end{bmatrix} \right) + b_z \end{aligned}$$


Minimize  $\sum_i (z_i - f(x_i, y_i))^2$ ,  
with respect to  $\mathbf{W}_1, \mathbf{w}_2, b_x, b_y, b_z$ .




# INTERPOLATING A SURFACE



# INTERPOLATING A SURFACE




$$z = 100 * (y - x^2)^2 + (1 - x)^2$$




4-node hidden layer


# ADDING MORE NODES




$$z = 100 * (y - x^2)^2 + (1 - x)^2$$



2 nodes -> loss 3.02e-01





3 nodes -> loss 2.08e-02




4 nodes -> loss 8.27e-03


# ADDING MORE NODES



2 nodes -> loss 2.61e-01




3 nodes -> loss 2.51e-04




4 nodes -> loss 3.07e-07




# MORE COMPLEX SURFACE




$$I = f(x, y)$$




# ADDING MORE NODES




$$I = f(x, y)$$




50 nodes -> loss 3.65e-01



100 nodes -> loss 2.50e-01



125 nodes -> loss 2.40e-01



300 nodes -> loss 1.92e-01



# UNIVERSAL APPROXIMATION THEOREM




A feedforward network with a linear output layer and at least one hidden layer with any 'squashing' activation function (e.g. logistic sigmoid) can approximate any Borel measurable function (from one finite-dimensional space to another) with any desired nonzero error.

Any continuous function on a closed and bounded set of  $\mathbb{R}^n$  is Borel-measurable.

—> In theory, any reasonable function can be approximated by a two-layer network as long as it is continuous.


# IN PRACTICE



- It may take an exponentially large number of parameters for a good approximation.
- The optimization problem becomes increasingly difficult.

—> The one hidden layer perceptron may not converge to the best solution!

# DEEP LEARNING




- The descriptive power of the net increases with the number of layers.
- In the case of a 1D signal, it is roughly proportional to  $\prod_n W_n$  where  $W_n$  represents the width of a layer.

# TWO HYPERPLANES

## ONE SINGLE LAYER

$$\begin{cases} \mathbf{h} = \max(\mathbf{W}\mathbf{x} + \mathbf{b}, 0) \\ y = \mathbf{w}'^\top \mathbf{h} \end{cases}$$

with  $\dim(\mathbf{h}) = 2$



# TWO HYPERPLANES

## TWO LAYERS

$$\begin{cases} \mathbf{h} = \max(\mathbf{W}\mathbf{x} + \mathbf{b}, 0) \\ \mathbf{h}' = \max(\mathbf{W}'\mathbf{h} + \mathbf{b}', 0) \\ y = \mathbf{w}''^\top \mathbf{h}' \end{cases}$$


$$\mathbf{h} = \begin{bmatrix} 0 \\ \mathbf{W}_{2,:}\mathbf{x} + b_2 \end{bmatrix}$$

$$\mathbf{h}' = \begin{bmatrix} \max(W'_{12}(\mathbf{W}_{2,:}\mathbf{x} + b_2) + b'_1, 0) \\ \max(W'_{22}(\mathbf{W}_{2,:}\mathbf{x} + b_2) + b'_2, 0) \end{bmatrix} = \begin{bmatrix} \max((W'_{12}\mathbf{W}_{2,:})\mathbf{x} + (b_2 + b'_1), 0) \\ \max((W'_{22}\mathbf{W}_{2,:})\mathbf{x} + (b_2 + b'_2), 0) \end{bmatrix}$$

$$\mathbf{h} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$\mathbf{h}' = \max(\mathbf{b}', 0)$$


$$y = \mathbf{w}'' \max(\mathbf{b}', 0) = \text{cst}$$



# BACK TO HYPERPLANES




The function learned by a DNN with reLu is

- piecewise affine;
- continuous;
- with regions related in a complex way.





# ADDING A SECOND LAYER






$$I = f(x, y)$$

1 Layer: 125 nodes -> loss 2.40e-01    2 Layers: 20 nodes -> loss 8.31e-02

501 weights in both cases

# ADDING A THIRD LAYER



$$I = f(x, y)$$




2 Layers: 20 nodes -> loss 8.31e-02

501 weights

3 Layers: 14 nodes -> loss 7.55e-02

477 weights

# ADDING A THIRD LAYER



$$I = f(x, y)$$



3 Layers: 15 nodes -> loss 5.93e-02

541 weights

3 Layers: 19 nodes -> loss 4.38e-02

837 weights


# MULTILAYER PERCEPTRONS




- Adding layers often yields better convergence properties.
- In current practice, deeper is usually better.




# SIMPLER WAY TO INTERPOLATE



Original 51x51 image:  
2601 gray level values.




Scaled 24x24 image:  
576 gray level values.



MLP 10/20/10 Interpolation:  
471 weights.

Simpler but not necessarily better!


# MLP TO RESNET




Further improvements in the convergence properties have been obtained by adding a bypass, which allows the final layers to only compute residuals.




# IMPROVING THE NETWORK





Original 51x51 image:  
2601 gray level values.




MLP 10/20/10 Interpolation:  
471 weights, loss  $6.43 \times 10^{-2}$ .



MLP 10/20/10/10 Interpolation:  
581 weights, loss  $5.30 \times 10^{-2}$ .




# IMPROVING THE NETWORK




- Relatively small improvement in **this** case.
- The problem is probably too small.  
—> Networks can behave very differently for small and large problems!


# TANH vs ReLU




Original 51x51 image:  
2601 gray level values.



MLP 10/20/10 Interpolation:  
**Tanh**, loss  $6.43e-02$ .




MLP 10/20/10 Interpolation:  
**ReLU**, loss  $3.07e-1$ .



- Tanh, 1 layers
- Tanh, 2 layers
- Tanh, 3 layers
- ReLU, 3 layers
- Tanh, 4 layers

# TANH vs ReLU



- Tanh works better than ReLU in **this** case.
- ReLU is widely credited with eliminating the vanishing gradient problem in large networks.

—> There is no substitute for experimentation!

# MULTILAYER PERCEPTRONS





The function learned by a DNN using either the ReLU or Tanh operators is:

- piecewise affine or smooth;
- continuous because it is a composition of continuous functions.

Each region created by a layer is split into smaller regions:

- The equations for each one are correlated in a complex way.
- This may explain why deeper networks generalize better than larger networks for a given number of parameters.

# STRENGTHS AND LIMITATIONS



- Powerful regressors but require many parameters, and therefore large training databases.
- Excellent at interpolation but less good at extrapolation. The training data must cover all cases of interest.


# IMAGE SPECIFICITIES



- In a typical image, the values of neighboring pixels tend to be more highly correlated than those of distant ones.
  - An image filter should be translation invariant.
- > These two properties can be exploited to drastically reduce the number of weights required by CNNs using so-called convolutional layers.




# CONVOLUTIONAL LAYER



$$\sigma \left( b + \sum_{x=0}^{n_x} \sum_{y=0}^{n_y} w_{i,j} a_{i+x,j+y} \right)$$




# FEATURE MAPS




# POOLING LAYER

hidden neurons (output from feature map)




max-pooling units




- Reduce the number of inputs by replacing all activations in a neighborhood by a single one.
- Can be thought as asking if a particular feature is present in that neighborhood while ignoring the exact location.

# ADDING THE POOLING LAYERS




The output size is reduced by the pooling layers.

# ADDING A FULLY CONNECTED LAYER



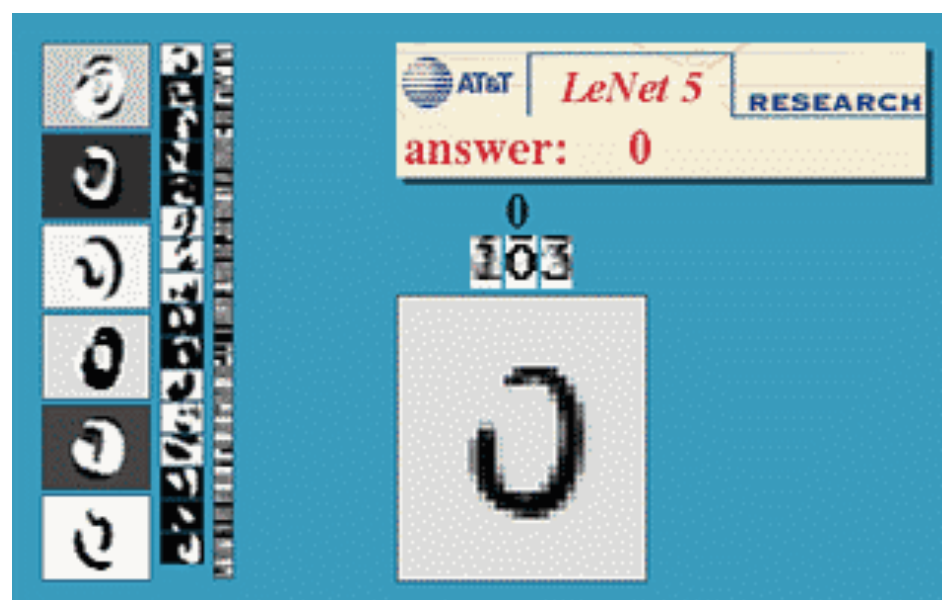

- Each neuron in the final fully connected layer is connected to all neurons in the preceding one.
- Deep architecture with many parameters to learn but still far fewer than an equivalent multilayer perceptron.

# MNIST




- The network takes as input 28x28 images represented as 784D vectors.
- The output is a 10D vector giving the probability of the image representing any of the 10 digits.
- There are 50'000 training pairs of images and the corresponding label, 10'000 validation pairs, and 5'000 testing pairs.

# LeNet (1989-1999)




# IS MAX POOL REQUIRED?



Accuracy	Train	Test
Conv 5x5, stride 1 Max pool 2x3	99.58	98.77
Conv 5x5, stride 2	99.42	98.31
Conv 5x5, stride 1 Conv 3x3, stride 2	99.38	98.57

# AlexNet (2012)



Task: Image classification


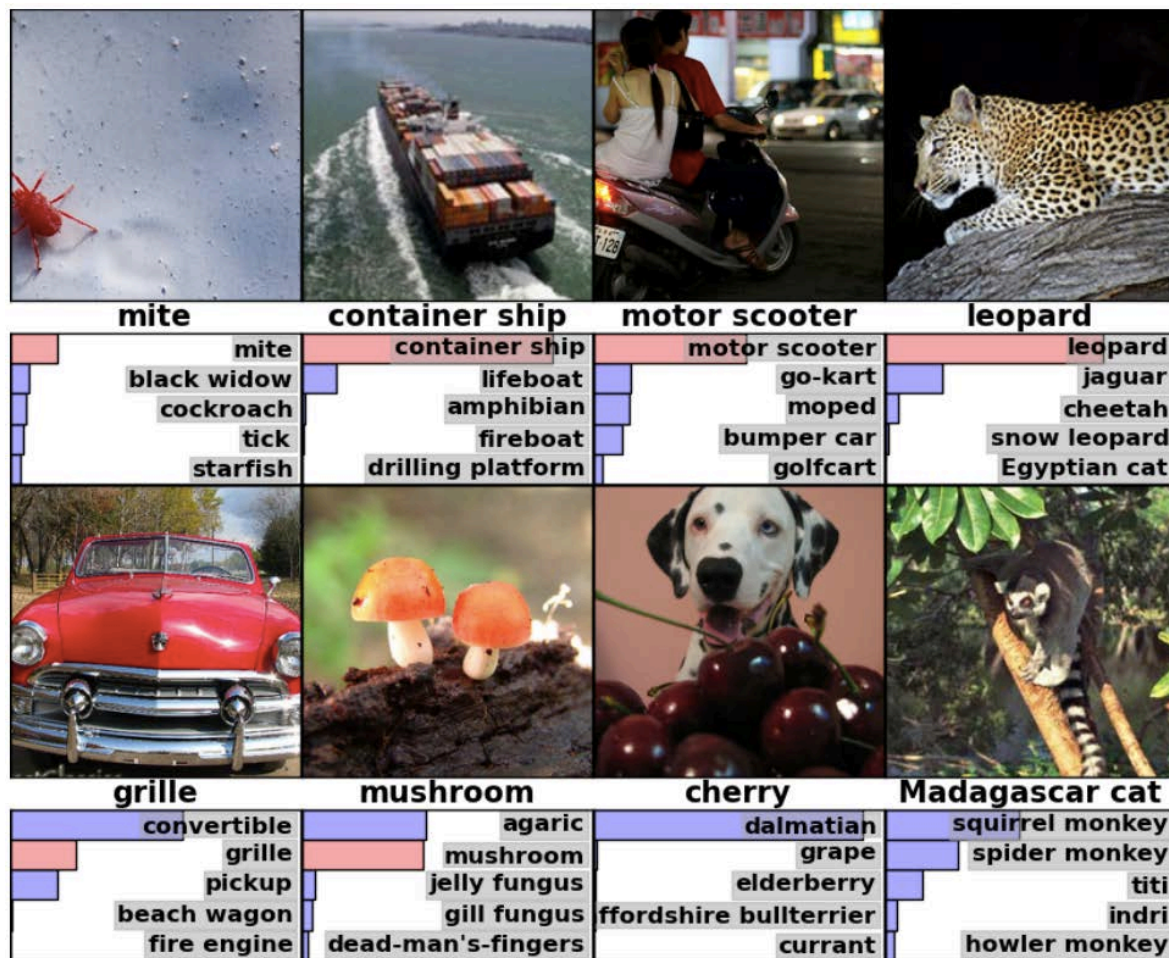
Training images: Large Scale Visual Recognition Challenge 2010

Training time: 2 weeks on 2 GPUs

Major Breakthrough: Training large networks has now been shown to be practical!!





# AlexNet RESULTS



- At the 2012 ImageNet Large Scale Visual Recognition Challenge, AlexNet achieved a top-5 error of 15.3%, more than 10.8% lower than the runner up.
- Since 2015, networks outperform humans on this task.

# FEATURE MAPS





First convolutional layer

Second convolutional layer



- Some of the convolutional masks seem very similar to oriented Gaussian or Gabor filters!
- Much ongoing work to better understand this.

# HIGHER ORDER DERIVATIVES






# FILTER BANKS




Hand-Designed



Learned

# SIZE AND DEPTH MATTER



"hibiscus"




"dahlia"

VGG19, 3 weeks of training.

GoogleLeNet.


“It was demonstrated that the representation depth is beneficial for the classification accuracy, and that state-of-the-art performance on the ImageNet challenge dataset can be achieved using a conventional ConvNet architecture.”

# HAND POSE ESTIMATION (2015)




Input: Depth image.

Output: 3D pose vector.





# REGRESSION




Network parameters are found by minimizing an objective function of the form

$$\min_{\mathbf{W}_l, \mathbf{B}_l} \sum_i ||\mathbf{F}(\mathbf{x}_i, \mathbf{W}_1, \dots, \mathbf{W}_L, \mathbf{b}_1, \dots, \mathbf{b}_L) - \mathbf{y}_i||^2$$


using

- stochastic gradient descent on mini-batches,
- dropout,
- hard example mining,
- .....

# ROC HUNTING



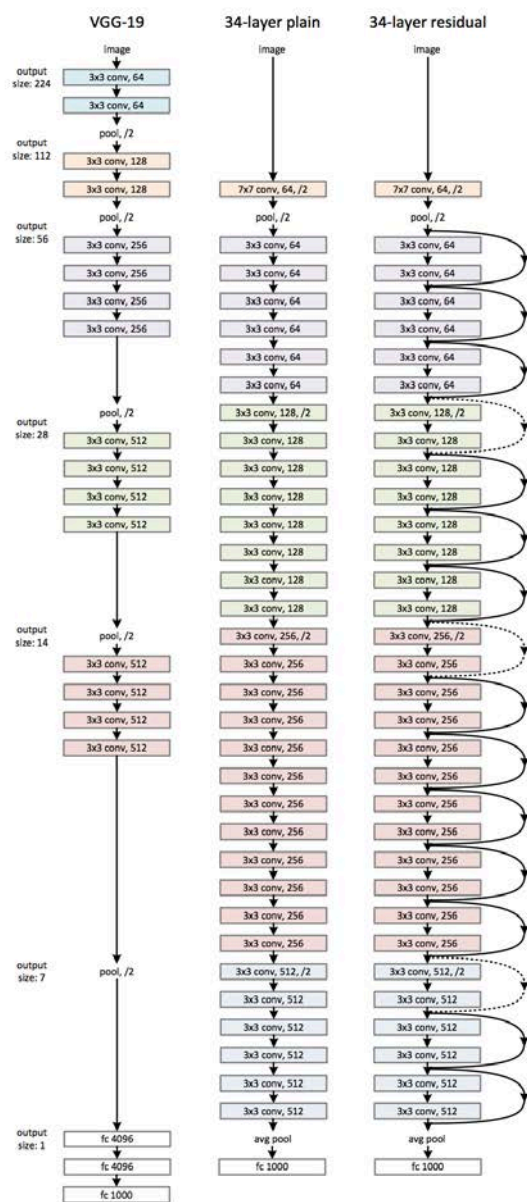


DeepFace  
Taigman et al. 2014

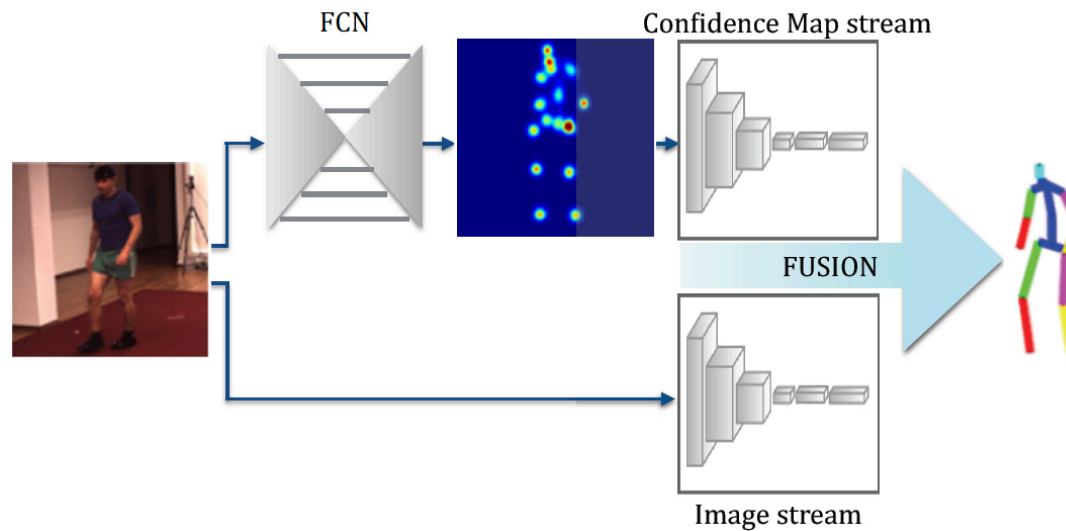





Deep Edge Detection  
Shen et al. 2015

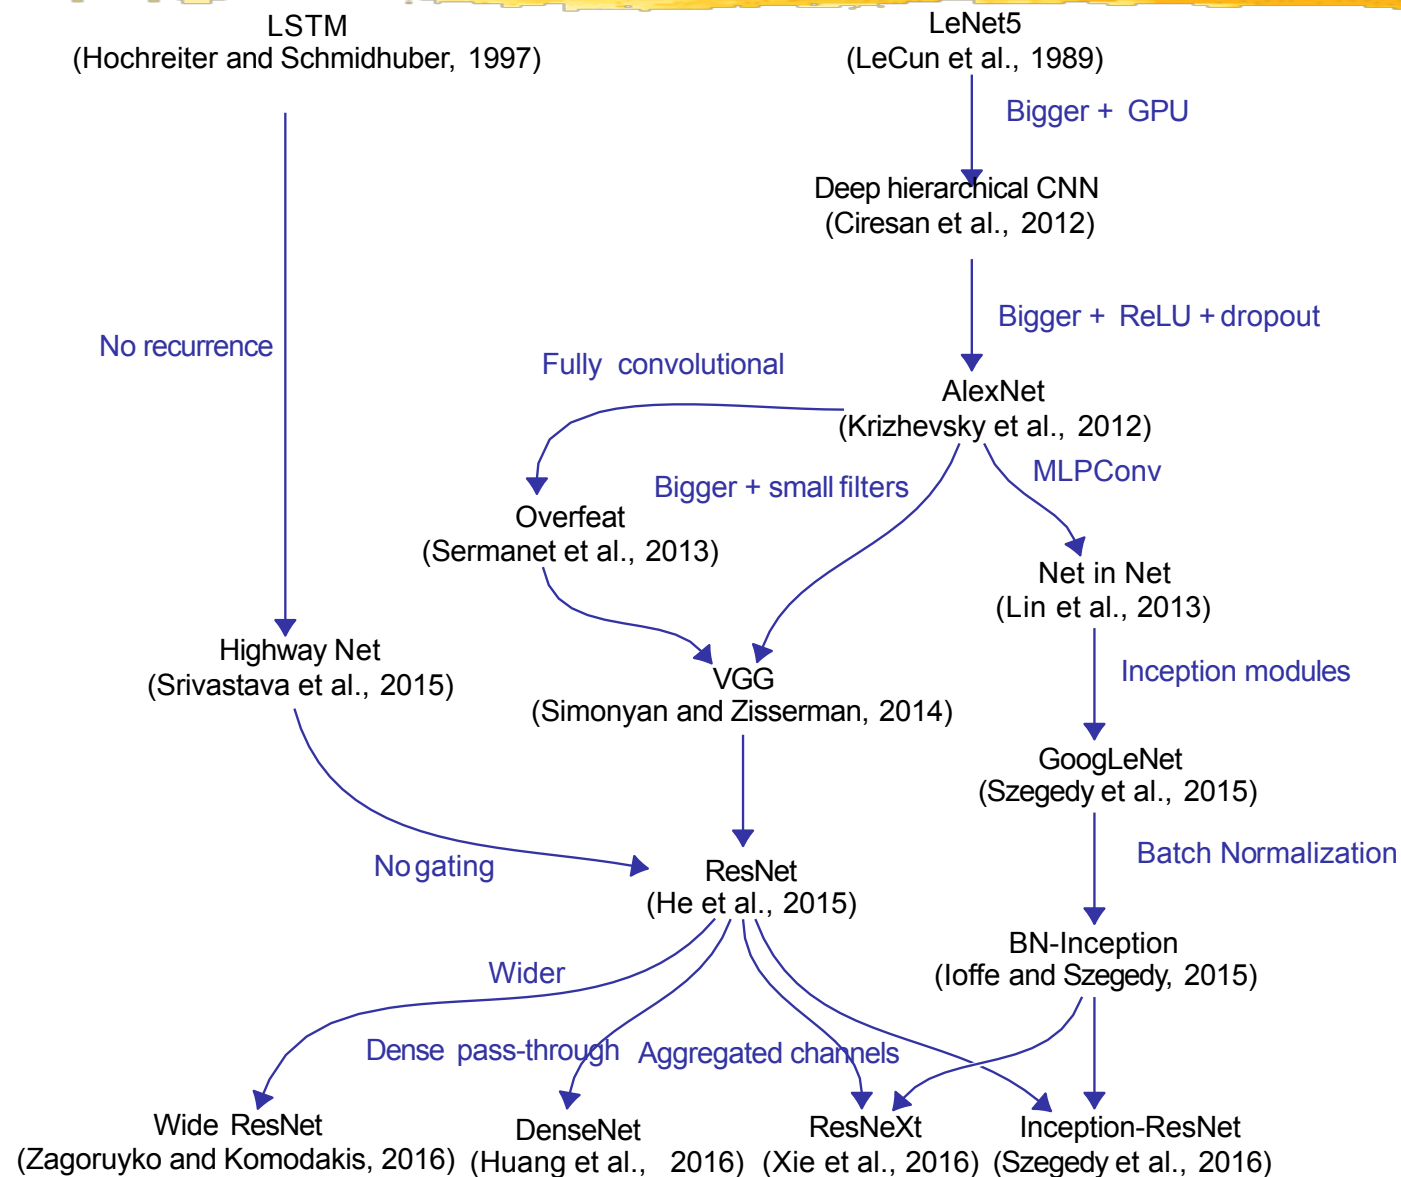
# DEEPER AND DEEPER



# MONOCULAR POSE ESTIMATION



# IMAGE CLASSIFICATION TAXONOMY



# ANOTHER POINT OF VIEW





To summarize roughly the evolution of convnets for image classification:

- Standard ones are extensions of LeNet5.
- Everybody loves ReLU.
- Newer ones have 100s of channels and 10s of layers.
- They can (should?) be fully convolutional.
- Pass-through connections allow deeper “residual” nets.
- Bottleneck local structures reduce the number of parameters.
- Aggregated pathways reduce the number of parameters.



# VISUAL CORTEX




# AlphaGo



- Uses Deep Nets to find the most promising locations to focus on.
- Performs Tree based search when possible.
- Relies on reinforcement learning and other ML techniques to train.


# ADVERSARIAL IMAGES



Szegedy et al. 2013



# XKCD'S VIEW ON THE MATTER



# IN SHORT



- Deep Belief Networks in general and Convolutional Neural Nets in particular outperform conventional Computer Vision algorithms on many benchmarks.
  - It is not fully understood why and unexpected failure cases have been demonstrated.
  - They require a lot of manual tuning to perform well and performance is hard to predict.
- > Many questions are still open and there is much work left to do.