

COM303: Digital Signal Processing

Lecture 21: Image Compression

- ▶ introduction to quantization
- ▶ the problem of image compression
- ▶ the JPEG standard

quantization: the basics

Quantization

- ▶ digital devices can only deal with integers (R bits per sample)
- ▶ we need to map the range of a signal onto a finite set of values
- ▶ irreversible loss of information \rightarrow quantization noise


Quantization

- ▶ digital devices can only deal with integers (R bits per sample)
- ▶ we need to map the range of a signal onto a finite set of values
- ▶ irreversible loss of information → quantization noise

Quantization

- ▶ digital devices can only deal with integers (R bits per sample)
- ▶ we need to map the range of a signal onto a finite set of values
- ▶ irreversible loss of information \rightarrow quantization noise


Quantization strategies



Several factors at play:

- ▶ storage budget (R bits per sample)
- ▶ storage scheme (fixed point, floating point)
- ▶ properties of the input
 - range
 - probability distribution


Quantization strategies



Several factors at play:

- ▶ storage budget (R bits per sample)
- ▶ storage scheme (fixed point, floating point)
- ▶ properties of the input
 - range
 - probability distribution


Quantization strategies



Several factors at play:

- ▶ storage budget (R bits per sample)
- ▶ storage scheme (fixed point, floating point)
- ▶ properties of the input
 - range
 - probability distribution


Quantization strategies



Several factors at play:

- ▶ storage budget (R bits per sample)
- ▶ storage scheme (fixed point, floating point)
- ▶ properties of the input
 - range
 - probability distribution


Quantization strategies



Several factors at play:

- ▶ storage budget (R bits per sample)
- ▶ storage scheme (fixed point, floating point)
- ▶ properties of the input
 - range
 - probability distribution


Uniform scalar quantization



The simplest quantizer:

- ▶ each sample is encoded individually (hence *scalar*) using R bits
- ▶ each sample is quantized independently (memoryless quantization)
- ▶ the input range is divided into 2^R equal-size intervals


Uniform scalar quantization



The simplest quantizer:

- ▶ each sample is encoded individually (hence *scalar*) using R bits
- ▶ each sample is quantized independently (memoryless quantization)
- ▶ the input range is divided into 2^R equal-size intervals

Uniform scalar quantization




The simplest quantizer:


- ▶ each sample is encoded individually (hence *scalar*) using R bits
- ▶ each sample is quantized independently (memoryless quantization)
- ▶ the input range is divided into 2^R equal-size intervals

Uniform scalar quantization

- ▶ input assumed uniformly distributed over $[A, B]$
- ▶ range is split into 2^R *equal* intervals of width $\Delta = (B - A)2^{-R}$
- ▶ quantized value is interval's midpoint



Uniform 3-Bit quantization function



Quantization Error

$$e[n] = \mathcal{Q}\{x[n]\} - x[n] = \hat{x}[n] - x[n]$$

- ▶ model $x[n]$ as a stochastic process
- ▶ model error as a white noise sequence:
 - error samples are uncorrelated
 - all error samples have the same distribution

Quantization Error

$$e[n] = \mathcal{Q}\{x[n]\} - x[n] = \hat{x}[n] - x[n]$$

- ▶ model $x[n]$ as a stochastic process
- ▶ model error as a white noise sequence:
 - error samples are uncorrelated
 - all error samples have the same distribution

Quantization Error

$$e[n] = \mathcal{Q}\{x[n]\} - x[n] = \hat{x}[n] - x[n]$$

- ▶ model $x[n]$ as a stochastic process
- ▶ model error as a white noise sequence:
 - error samples are uncorrelated
 - all error samples have the same distribution

Quantization Error

$$e[n] = \mathcal{Q}\{x[n]\} - x[n] = \hat{x}[n] - x[n]$$

- ▶ model $x[n]$ as a stochastic process
- ▶ model error as a white noise sequence:
 - error samples are uncorrelated
 - all error samples have the same distribution

Error analysis for uniform iid input

- ▶ error energy

$$\sigma_e^2 = \Delta^2/12, \quad \Delta = (B - A)/2^R$$

- ▶ signal energy

$$\sigma_x^2 = (B - A)^2/12$$

- ▶ signal to noise ratio

$$\text{SNR} = 2^{2R}$$

- ▶ in dB

$$\text{SNR}_{\text{dB}} = 10 \log_{10} 2^{2R} \approx 6R \text{ dB}$$

Error analysis for uniform iid input

- ▶ error energy

$$\sigma_e^2 = \Delta^2/12, \quad \Delta = (B - A)/2^R$$

- ▶ signal energy

$$\sigma_x^2 = (B - A)^2/12$$

- ▶ signal to noise ratio

$$\text{SNR} = 2^{2R}$$

- ▶ in dB

$$\text{SNR}_{\text{dB}} = 10 \log_{10} 2^{2R} \approx 6R \text{ dB}$$

Error analysis for uniform iid input

- ▶ error energy

$$\sigma_e^2 = \Delta^2/12, \quad \Delta = (B - A)/2^R$$

- ▶ signal energy

$$\sigma_x^2 = (B - A)^2/12$$

- ▶ signal to noise ratio

$$\text{SNR} = 2^{2R}$$

- ▶ in dB

$$\text{SNR}_{\text{dB}} = 10 \log_{10} 2^{2R} \approx 6R \text{ dB}$$

Error analysis for uniform iid input

- ▶ error energy

$$\sigma_e^2 = \Delta^2/12, \quad \Delta = (B - A)/2^R$$

- ▶ signal energy

$$\sigma_x^2 = (B - A)^2/12$$

- ▶ signal to noise ratio

$$\text{SNR} = 2^{2R}$$

- ▶ in dB

$$\text{SNR}_{\text{dB}} = 10 \log_{10} 2^{2R} \approx 6R \text{ dB}$$

The “6dB/bit” rule of thumb

- ▶ a compact disk has 16 bits/sample:

$$\max \text{SNR} = 96\text{dB}$$

- ▶ a DVD has 24 bits/sample:

$$\max \text{SNR} = 144\text{dB}$$

The “6dB/bit” rule of thumb

- ▶ a compact disk has 16 bits/sample:

$$\max \text{ SNR} = 96\text{dB}$$

- ▶ a DVD has 24 bits/sample:

$$\max \text{ SNR} = 144\text{dB}$$

Rate/Distortion Curve




image compression fundamentals

A thought experiment

- ▶ consider all possible 256×256 , 8bpp images
- ▶ each image is 524,288 bits
- ▶ total number of possible images: $2^{524,288} \approx 10^{157,826}$
- ▶ number of atoms in the universe: 10^{82}

A thought experiment

- ▶ consider all possible 256×256 , 8bpp images
- ▶ each image is 524,288 bits
- ▶ total number of possible images: $2^{524,288} \approx 10^{157,826}$
- ▶ number of atoms in the universe: 10^{82}

A thought experiment

- ▶ consider all possible 256×256 , 8bpp images
- ▶ each image is 524,288 bits
- ▶ total number of possible images: $2^{524,288} \approx 10^{157,826}$
- ▶ number of atoms in the universe: 10^{82}

A thought experiment

- ▶ consider all possible 256×256 , 8bpp images
- ▶ each image is 524,288 bits
- ▶ total number of possible images: $2^{524,288} \approx 10^{157,826}$
- ▶ number of atoms in the universe: 10^{82}

How many bits per image?

Another thought experiment

- ▶ take *all* images in the world and list them in an “encyclopedia of images”
- ▶ to indicate an image, simply give its number
- ▶ on the Internet: $M = 50$ billion
- ▶ raw encoding: 524,288 bits per image
- ▶ enumeration-based encoding: $\log_2 M \approx 36$ bits per image
- ▶ (of course, side information is HUGE)

How many bits per image?

Another thought experiment

- ▶ take *all* images in the world and list them in an “encyclopedia of images”
- ▶ to indicate an image, simply give its number
- ▶ on the Internet: $M = 50$ billion
- ▶ raw encoding: 524,288 bits per image
- ▶ enumeration-based encoding: $\log_2 M \approx 36$ bits per image
- ▶ (of course, side information is HUGE)

How many bits per image?

Another thought experiment

- ▶ take *all* images in the world and list them in an “encyclopedia of images”
- ▶ to indicate an image, simply give its number
- ▶ on the Internet: $M = 50$ billion
- ▶ raw encoding: 524,288 bits per image
- ▶ enumeration-based encoding: $\log_2 M \approx 36$ bits per image
- ▶ (of course, side information is HUGE)

How many bits per image?

Another thought experiment

- ▶ take *all* images in the world and list them in an “encyclopedia of images”
- ▶ to indicate an image, simply give its number
- ▶ on the Internet: $M = 50$ billion
- ▶ raw encoding: 524,288 bits per image
- ▶ enumeration-based encoding: $\log_2 M \approx 36$ bits per image
- ▶ (of course, side information is HUGE)

How many bits per image?

Another thought experiment

- ▶ take *all* images in the world and list them in an “encyclopedia of images”
- ▶ to indicate an image, simply give its number
- ▶ on the Internet: $M = 50$ billion
- ▶ raw encoding: 524,288 bits per image
- ▶ enumeration-based encoding: $\log_2 M \approx 36$ bits per image
- ▶ (of course, side information is HUGE)

How many bits per image?

Another thought experiment

- ▶ take *all* images in the world and list them in an “encyclopedia of images”
- ▶ to indicate an image, simply give its number
- ▶ on the Internet: $M = 50$ billion
- ▶ raw encoding: 524,288 bits per image
- ▶ enumeration-based encoding: $\log_2 M \approx 36$ bits per image
- ▶ (of course, side information is HUGE)

Compression

Another approach:

- ▶ exploit “physical” redundancy
- ▶ allocate bits for things that matter (e.g. edges)
- ▶ use psychovisual experiments to find out what matters
- ▶ some information is discarded: *lossy* compression

Compression

Another approach:

- ▶ exploit “physical” redundancy
- ▶ allocate bits for things that matter (e.g. edges)
- ▶ use psychovisual experiments to find out what matters
- ▶ some information is discarded: *lossy* compression

Compression

Another approach:

- ▶ exploit “physical” redundancy
- ▶ allocate bits for things that matter (e.g. edges)
- ▶ use psychovisual experiments to find out what matters
- ▶ some information is discarded: *lossy* compression

Compression

Another approach:

- ▶ exploit “physical” redundancy
- ▶ allocate bits for things that matter (e.g. edges)
- ▶ use psychovisual experiments to find out what matters
- ▶ some information is discarded: *lossy* compression

Key ingredients

- ▶ compressing at block level
- ▶ using a suitable transform (i.e., a change of basis)
- ▶ smart quantization
- ▶ entropy coding

Key ingredients

- ▶ compressing at block level
- ▶ using a suitable transform (i.e., a change of basis)
- ▶ smart quantization
- ▶ entropy coding

Key ingredients

- ▶ compressing at block level
- ▶ using a suitable transform (i.e., a change of basis)
- ▶ smart quantization
- ▶ entropy coding

Key ingredients

- ▶ compressing at block level
- ▶ using a suitable transform (i.e., a change of basis)
- ▶ smart quantization
- ▶ entropy coding

Compressing at pixel level

- ▶ reduce number bits per pixel
- ▶ equivalent to coarser quantization
- ▶ in the limit, 1bpp

Compressing at pixel level

- ▶ reduce number bits per pixel
- ▶ equivalent to coarser quantization
- ▶ in the limit, 1bpp

Compressing at pixel level

- ▶ reduce number bits per pixel
- ▶ equivalent to coarser quantization
- ▶ in the limit, 1bpp



Compressing at block level

- ▶ divide the image in blocks
- ▶ code the average value with 8 bits
- ▶ 3×3 blocks at 8 bits per block gives less than 1bpp

Compressing at block level

- ▶ divide the image in blocks
- ▶ code the average value with 8 bits
- ▶ 3×3 blocks at 8 bits per block gives less than 1bpp

Compressing at block level

- ▶ divide the image in blocks
- ▶ code the average value with 8 bits
- ▶ 3×3 blocks at 8 bits per block gives less than 1bpp



Compressing at block level

- ▶ exploit the local spatial correlation
- ▶ compress remote regions independently


Compressing at block level

- ▶ exploit the local spatial correlation
- ▶ compress remote regions independently

Transform coding

A simple example:


- ▶ take a DT signal, assume R bits per sample
- ▶ storing the signal requires NR bits
- ▶ now you take the DFT and it looks like this
- ▶ in theory, we can just code the two nonzero DFT coefficients!



Transform coding

A simple example:


- ▶ take a DT signal, assume R bits per sample
- ▶ storing the signal requires NR bits
- ▶ now you take the DFT and it looks like this
- ▶ in theory, we can just code the two nonzero DFT coefficients!



Transform coding

A simple example:


- ▶ take a DT signal, assume R bits per sample
- ▶ storing the signal requires NR bits
- ▶ now you take the DFT and it looks like this
- ▶ in theory, we can just code the two nonzero DFT coefficients!



Transform coding

A simple example:

- ▶ take a DT signal, assume R bits per sample
- ▶ storing the signal requires NR bits
- ▶ now you take the DFT and it looks like this
- ▶ in theory, we can just code the two nonzero DFT coefficients!



Transform coding

Ideally, we would like a transform that:

- ▶ captures the important features of an image block in a few coefficients
- ▶ is efficient to compute
- ▶ answer: the Discrete Cosine Transform

Transform coding

Ideally, we would like a transform that:

- ▶ captures the important features of an image block in a few coefficients
- ▶ is efficient to compute
- ▶ answer: the Discrete Cosine Transform


Transform coding

Ideally, we would like a transform that:

- ▶ captures the important features of an image block in a few coefficients
- ▶ is efficient to compute
- ▶ answer: the Discrete Cosine Transform

$$C[k_1, k_2] = \sum_{n_1=0}^{N-1} \sum_{n_2=0}^{N-1} x[n_1, n_2] \cos \left[\frac{\pi}{N} \left(n_1 + \frac{1}{2} \right) k_1 \right] \cos \left[\frac{\pi}{N} \left(n_2 + \frac{1}{2} \right) k_2 \right]$$

DCT basis vectors for an 8×8 image



Smart quantization

- ▶ deadzone
- ▶ variable step (fine to coarse)


Smart quantization

- ▶ deadzone
- ▶ variable step (fine to coarse)

Quantization

Standard quantization:


$$\hat{x} = \text{floor}(x) + 0.5$$



Quantization

Deadzone quantization:

$$\hat{x} = \text{round}(x)$$



Entropy coding

- ▶ minimize the effort to encode a certain amount of information
- ▶ associate short symbols to frequent values and vice-versa
- ▶ if it sounds familiar it's because it is...

Entropy coding

- ▶ minimize the effort to encode a certain amount of information
- ▶ associate short symbols to frequent values and vice-versa
- ▶ if it sounds familiar it's because it is...

Entropy coding

- ▶ minimize the effort to encode a certain amount of information
- ▶ associate short symbols to frequent values and vice-versa
- ▶ if it sounds familiar it's because it is...

Entropy coding

A ● ■
B ■ ● ● ●
C ■ ● ■ ●
D ■ ● ●
E ●
F ● ● ■ ●
G ■ ■ ●
H ● ● ● ●
I ● ●
J ● ■ ■ ■ ■
K ■ ● ■
L ● ■ ● ●
M ■ ■
N ■ ●
O ■ ■ ■
P ● ■ ■ ■ ●
Q ■ ■ ■ ● ■
R ● ■ ■ ●
S ● ● ●
T ■

U ● ● ■
V ● ● ■ ■
W ● ■ ■
X ■ ● ● ■
Y ■ ● ■ ■ ■
Z ■ ■ ● ●

1 ● ■ ■ ■ ■ ■
2 ● ● ■ ■ ■ ■
3 ● ● ● ■ ■ ■
4 ● ● ● ● ■
5 ● ● ● ● ●
6 ■ ● ● ● ●
7 ■ ■ ■ ● ● ●
8 ■ ■ ■ ■ ● ●
9 ■ ■ ■ ■ ■ ●
0 ■ ■ ■ ■ ■ ■

the JPEG standard

Key ingredients

- ▶ compressing at block level
- ▶ using a suitable transform (i.e., a change of basis)
- ▶ smart quantization
- ▶ entropy coding

Key ingredients

- ▶ split image into 8×8 non-overlapping blocks
- ▶ using a suitable transform (i.e., a change of basis)
- ▶ smart quantization
- ▶ entropy coding

Key ingredients

- ▶ split image into 8×8 non-overlapping blocks
- ▶ compute the DCT of each block
- ▶ smart quantization
- ▶ entropy coding


Key ingredients

- ▶ split image into 8×8 non-overlapping blocks
- ▶ compute the DCT of each block
- ▶ quantize DCT coefficients according to psychovisually-tuned tables
- ▶ entropy coding


Key ingredients

- ▶ split image into 8×8 non-overlapping blocks
- ▶ compute the DCT of each block
- ▶ quantize DCT coefficients according to psychovisually-tuned tables
- ▶ run-length encoding and Huffman coding

DCT coefficients of image blocks (detail)



DCT coefficients of image blocks (detail)



Smart quantization

- ▶ most coefficients are negligible \rightarrow captured by the deadzone
- ▶ some coefficients have a higher visual impact
- ▶ find out the critical coefficients by experimentation
- ▶ use smaller quantization intervals (i.e. use more more bits) for the important coefficients

Smart quantization

- ▶ most coefficients are negligible \rightarrow captured by the deadzone
- ▶ some coefficients have a higher visual impact
- ▶ find out the critical coefficients by experimentation
- ▶ use smaller quantization intervals (i.e. use more more bits) for the important coefficients

Smart quantization

- ▶ most coefficients are negligible \rightarrow captured by the deadzone
- ▶ some coefficients have a higher visual impact
- ▶ find out the critical coefficients by experimentation
- ▶ use smaller quantization intervals (i.e. use more more bits) for the important coefficients

Smart quantization

- ▶ most coefficients are negligible \rightarrow captured by the deadzone
- ▶ some coefficients have a higher visual impact
- ▶ find out the critical coefficients by experimentation
- ▶ use smaller quantization intervals (i.e. use more more bits) for the important coefficients

Psychovisually-tuned quantization table

$$\hat{c}[k_1, k_2] = \text{round}(c[k_1, k_2]/Q[k_1, k_2])$$

$$Q = \begin{bmatrix} 16 & 11 & 10 & 16 & 24 & 40 & 51 & 61 \\ 12 & 12 & 14 & 19 & 26 & 58 & 60 & 55 \\ 14 & 13 & 16 & 24 & 40 & 57 & 69 & 56 \\ 14 & 17 & 22 & 29 & 51 & 87 & 80 & 62 \\ 18 & 22 & 37 & 56 & 68 & 109 & 103 & 77 \\ 24 & 35 & 55 & 64 & 81 & 104 & 113 & 92 \\ 49 & 64 & 78 & 87 & 103 & 121 & 120 & 101 \\ 72 & 92 & 95 & 98 & 112 & 100 & 103 & 99 \end{bmatrix}$$

Advantages of nonuniform bit allocation at 0.2bpp



uniform



tuned

Advantages of nonuniform bit allocation (detail)



uniform



tuned

Efficient coding

- ▶ most coefficients are small, decreasing with index
- ▶ use zigzag scan to maximize ordering
- ▶ quantization will create long series of zeros


Efficient coding

- ▶ most coefficients are small, decreasing with index
- ▶ use zigzag scan to maximize ordering
- ▶ quantization will create long series of zeros

Efficient coding

- ▶ most coefficients are small, decreasing with index
- ▶ use zigzag scan to maximize ordering
- ▶ quantization will create long series of zeros

Zigzag scan



Example

$$\begin{bmatrix} 100 & -60 & 0 & 6 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 13 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Example

$$\begin{bmatrix} 100 & -60 & 0 & 6 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 13 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

[illegible]

Runlength encoding

- ▶ the DC value is encoded differentially wrt previous block
- ▶ each nonzero AC value is encoded as the triple

$$[(r, s), c]$$

- r is the *runlength* i.e. the number of zeros before the current value
- s is the *category* i.e. the number of bits needed to encode the value
- c is the actual value
- $(0, 0)$ indicates that from now on it's only zeros (end of block)

Runlength encoding

- ▶ the DC value is encoded differentially wrt previous block
- ▶ each nonzero AC value is encoded as the triple

$$[(r, s), c]$$

- r is the *runlength* i.e. the number of zeros before the current value
- s is the *category* i.e. the number of bits needed to encode the value
- c is the actual value
- $(0, 0)$ indicates that from now on it's only zeros (end of block)

Runlength encoding

- ▶ the DC value is encoded differentially wrt previous block
- ▶ each nonzero AC value is encoded as the triple

$$[(r, s), c]$$

- r is the *runlength* i.e. the number of zeros before the current value
- s is the *category* i.e. the number of bits needed to encode the value
- c is the actual value
- $(0, 0)$ indicates that from now on it's only zeros (end of block)

Runlength encoding

- ▶ the DC value is encoded differentially wrt previous block
- ▶ each nonzero AC value is encoded as the triple

$$[(r, s), c]$$

- r is the *runlength* i.e. the number of zeros before the current value
- s is the *category* i.e. the number of bits needed to encode the value
- c is the actual value
- $(0, 0)$ indicates that from now on it's only zeros (end of block)

Runlength encoding

- ▶ the DC value is encoded differentially wrt previous block
- ▶ each nonzero AC value is encoded as the triple

$$[(r, s), c]$$

- r is the *runlength* i.e. the number of zeros before the current value
- s is the *category* i.e. the number of bits needed to encode the value
- c is the actual value
- $(0, 0)$ indicates that from now on it's only zeros (end of block)

Runlength encoding

- ▶ the DC value is encoded differentially wrt previous block
- ▶ each nonzero AC value is encoded as the triple

$$[(r, s), c]$$

- r is the *runlength* i.e. the number of zeros before the current value
- s is the *category* i.e. the number of bits needed to encode the value
- c is the actual value
- $(0, 0)$ indicates that from now on it's only zeros (end of block)

Example

$$\begin{bmatrix} 100 & -60 & 0 & 6 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 13 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

$[100], [(0, 6), -60], [(4, 3), 6], [(3, 4), 13], [(8, 1), -1], [(0, 0)]$

Example

$$\begin{bmatrix} 100 & -60 & 0 & 6 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 13 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

$[100]$, $[(0, 6), -60]$, $[(4, 3), 6]$, $[(3, 4), 13]$, $[(8, 1), -1]$, $[(0, 0)]$

The runlength-size pairs

- ▶ by design, $(r, s) \in \mathcal{A}$ with $|\mathcal{A}| = 256$
- ▶ in theory, 8 bits per pair
- ▶ some pairs are much more common than others!
- ▶ a lot of space can be saved by being smart

The runlength-size pairs

- ▶ by design, $(r, s) \in \mathcal{A}$ with $|\mathcal{A}| = 256$
- ▶ in theory, 8 bits per pair
- ▶ some pairs are much more common than others!
- ▶ a lot of space can be saved by being smart

The runlength-size pairs

- ▶ by design, $(r, s) \in \mathcal{A}$ with $|\mathcal{A}| = 256$
- ▶ in theory, 8 bits per pair
- ▶ some pairs are much more common than others!
- ▶ a lot of space can be saved by being smart

The runlength-size pairs

- ▶ by design, $(r, s) \in \mathcal{A}$ with $|\mathcal{A}| = 256$
- ▶ in theory, 8 bits per pair
- ▶ some pairs are much more common than others!
- ▶ a lot of space can be saved by being smart

Variable-length encoding

great idea: shorter binary sequences for common symbols

A	• ■	U	• • ■
B	■ ■ • •	V	• • • ■
C	■ ■ • ■ •	W	• ■ ■ ■
D	■ ■ • •	X	■ ■ • • ■
E	•	Y	■ ■ • ■ ■
F	• • ■ ■ •	Z	■ ■ ■ • •
G	■ ■ ■ •		
H	• • • •		
I	• •		
J	• ■ ■ ■ ■		
K	■ ■ • ■	1	• ■ ■ ■ ■ ■
L	• ■ ■ •	2	• • ■ ■ ■ ■
M	■ ■ ■	3	• • • ■ ■ ■
N	■ ■ •	4	• • • • ■ ■
O	■ ■ ■ ■	5	• • • • •
P	• ■ ■ ■ •	6	■ ■ • • •
Q	■ ■ ■ • ■	7	■ ■ ■ • •
R	• ■ ■ •	8	■ ■ ■ ■ • •
S	• • •	9	■ ■ ■ ■ ■ •
T	■	0	■ ■ ■ ■ ■ ■

Variable-length encoding

however: if symbols have different lengths, we must know how to parse them!

- ▶ in English, spaces separate words → extra symbol (wasteful)
- ▶ in Morse code, pauses separate letters and words (wasteful)
- ▶ can we do away with separators?

Variable-length encoding

however: if symbols have different lengths, we must know how to parse them!

- ▶ in English, spaces separate words → extra symbol (wasteful)
- ▶ in Morse code, pauses separate letters and words (wasteful)
- ▶ can we do away with separators?

Variable-length encoding

however: if symbols have different lengths, we must know how to parse them!

- ▶ in English, spaces separate words → extra symbol (wasteful)
- ▶ in Morse code, pauses separate letters and words (wasteful)
- ▶ can we do away with separators?

Prefix-free codes

- ▶ no valid sequence can be the beginning of another valid sequence
- ▶ can parse a bitstream sequentially with no look-ahead
- ▶ extremely easy to understand graphically...


Prefix-free codes

- ▶ no valid sequence can be the beginning of another valid sequence
- ▶ can parse a bitstream sequentially with no look-ahead
- ▶ extremely easy to understand graphically...

Prefix-free codes


- ▶ no valid sequence can be the beginning of another valid sequence
- ▶ can parse a bitstream sequentially with no look-ahead
- ▶ extremely easy to understand graphically...

Prefix-free code




001100110101100

Prefix-free code



001100110101100


Prefix-free code



001100110101100

A


Prefix-free code



001100110101100

AA


Prefix-free code



00**11**00110101100

AAB


Prefix-free code



001100110101100

AABA


Prefix-free code



001100110101100

AABAA


Prefix-free code



001100**1**0101100

AABAAB


Prefix-free code



001100110101100

AABAABA


Prefix-free code



001100110101100

AABAABAD

Prefix-free code



001100110101100

AABAABADC

Entropy coding

goal: minimize message length

- ▶ assign short sequences to more frequent symbols
- ▶ the Huffman algorithm builds the optimal code for a set of symbol probabilities
- ▶ in JPEG, you can use a “general-purpose” Huffman code or build your own (but then you pay a “side-information” price)

Entropy coding

goal: minimize message length

- ▶ assign short sequences to more frequent symbols
- ▶ the Huffman algorithm builds the optimal code for a set of symbol probabilities
- ▶ in JPEG, you can use a “general-purpose” Huffman code or build your own (but then you pay a “side-information” price)

Entropy coding

goal: minimize message length

- ▶ assign short sequences to more frequent symbols
- ▶ the Huffman algorithm builds the optimal code for a set of symbol probabilities
- ▶ in JPEG, you can use a “general-purpose” Huffman code or build your own (but then you pay a “side-information” price)

Example

► four symbols: A, B, C, D

► probability table:

$$p(A) = 0.38$$

$$p(B) = 0.32$$

$$p(C) = 0.1$$

$$p(D) = 0.2$$

Example

- ▶ four symbols: A, B, C, D
- ▶ probability table:


$$p(A) = 0.38$$

$$p(B) = 0.32$$

$$p(C) = 0.1$$


$$p(D) = 0.2$$

Building the Huffman code




$$p(A) = 0.38 \quad p(B) = 0.32 \quad p(C) = 0.1 \quad p(D) = 0.2$$

Building the Huffman code



$$p(A) = 0.38 \quad p(B) = 0.32 \quad p(C + D) = 0.3$$

Huffman Coding



$$p(A) = 0.38 \quad p(B + C + D) = 0.62$$

Conclusions

- ▶ JPEG is a very complex and comprehensive standard:
 - lossless, lossy
 - color, B&W
 - progressive encoding
 - HDR (12bpp) for medical imaging
- ▶ JPEG is VERY good:
 - compression factor of 10:1 virtually indistinguishable
 - rates of 1bpp for RGB images acceptable (25:1 compression ratio)
- ▶ other important compression schemes:
 - TIFF, JPEG2000
 - MPEG (MP3)