



# **Data Locality: Computer Hardware and the Memory Hierarchy**

---

Christoph Koch

*School of Computer & Communication Sciences, EPFL*

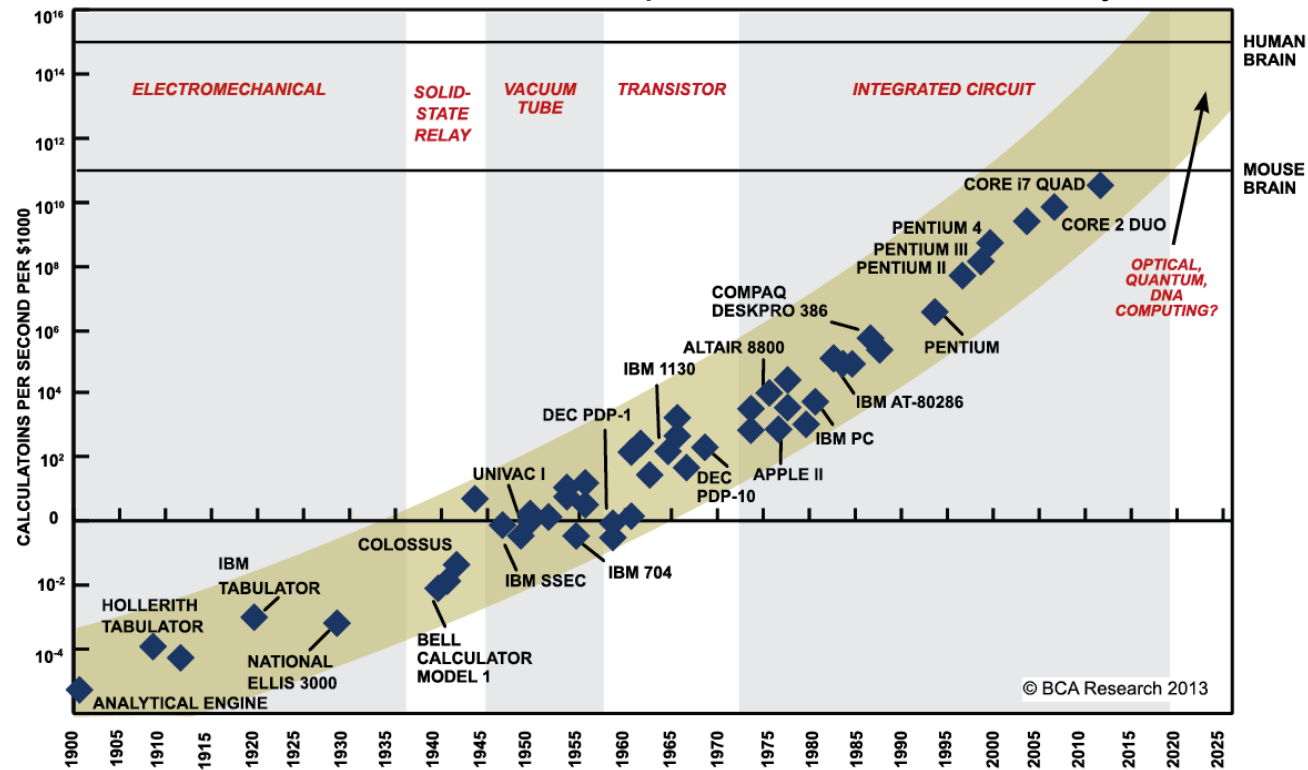
# Goals of this lecture

---

- Moore's law is failing:
  - *Computers don't get faster anymore.*
  - *We need to understand parallel computing, and its limitations.*
  - *We need to optimize our computations for locality.*
- Understanding data locality; principles of leveraging and maximizing locality.
- The memory hierarchy; caches.

# Moore's Law through the Classic Era

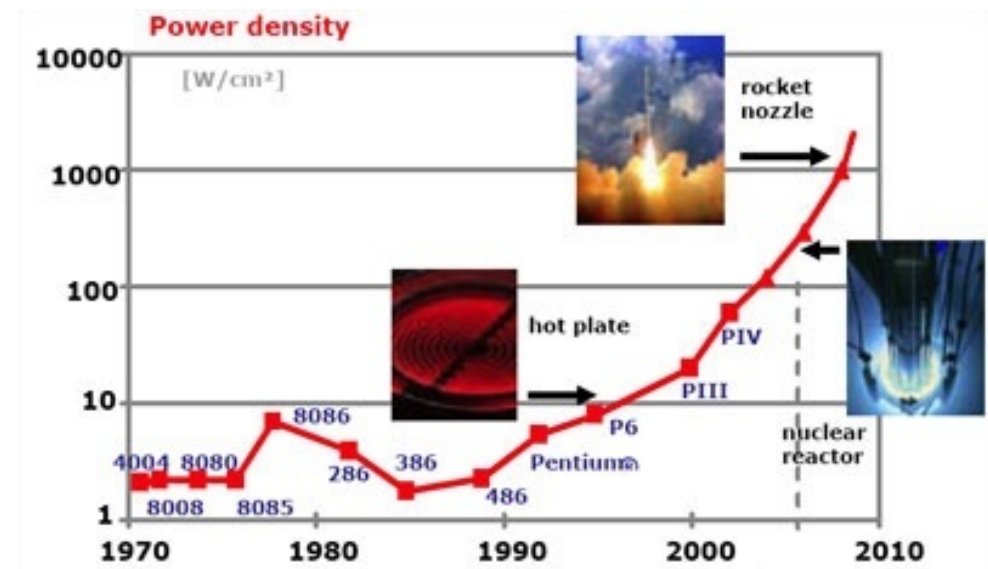
Moore's Law: #transistors achievable per  $\text{mm}^2$  doubles every ~18 months.



SOURCE: RAY KURZWEIL, "THE SINGULARITY IS NEAR: WHEN HUMANS TRANSCEND BIOLOGY", P.67, THE VIKING PRESS, 2006. DATAPPOINTS BETWEEN 2000 AND 2012 REPRESENT BCA ESTIMATES.

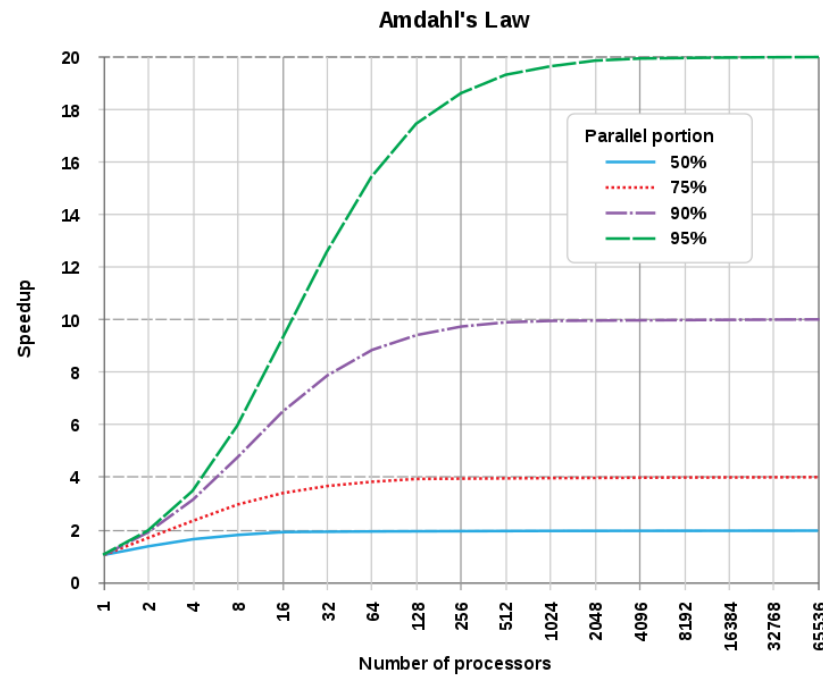
# Dennard Scaling has failed

- Dennard scaling: As integration increases (transistor size decreases) by 2x, required voltage decreases by 4x, so power density remains constant.
- Dennard scaling has failed
  - Due to quantum effects
- Consequence: we can't shrink logic/ increase clock rates anymore and still cool the chips.



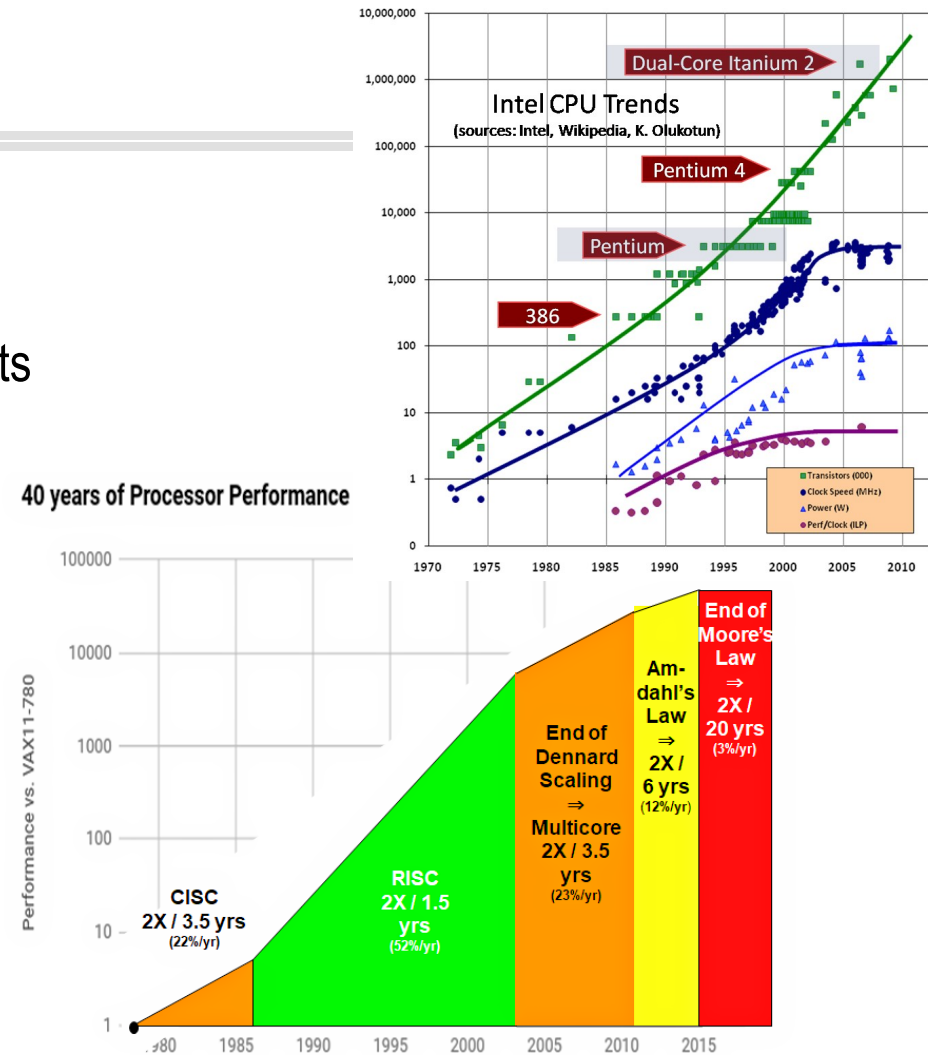
# Amdahl's law

The benefits of parallelization are limited due to the latency caused by the inherently sequential part of the computation.



# Moore's Law is Failing

- Failure of Dennard Scaling
- Amdahl's law – parallelization has its limits
- Can't communicate a bit/signal with less than one electron (or anything approaching it) – can't shrink transistors further after ~2030.
- Consequence: computation becomes less local.



# Locality

---

Locality relates (software) systems with the physical world.

- We can fundamentally only pack so much memory and computing power into a limited space.
- As data and our computing challenges grow, we get into locality problems.
- We can reduce pretty much any performance concern to a locality issue.

# The role of locality

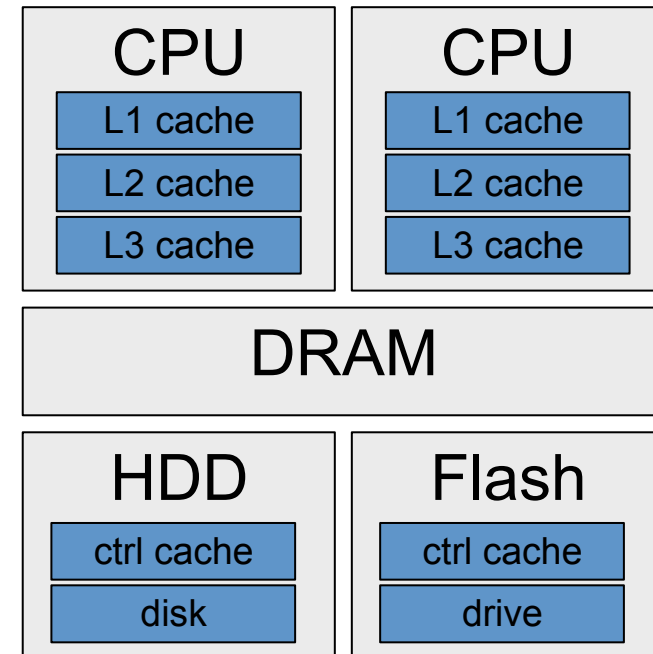
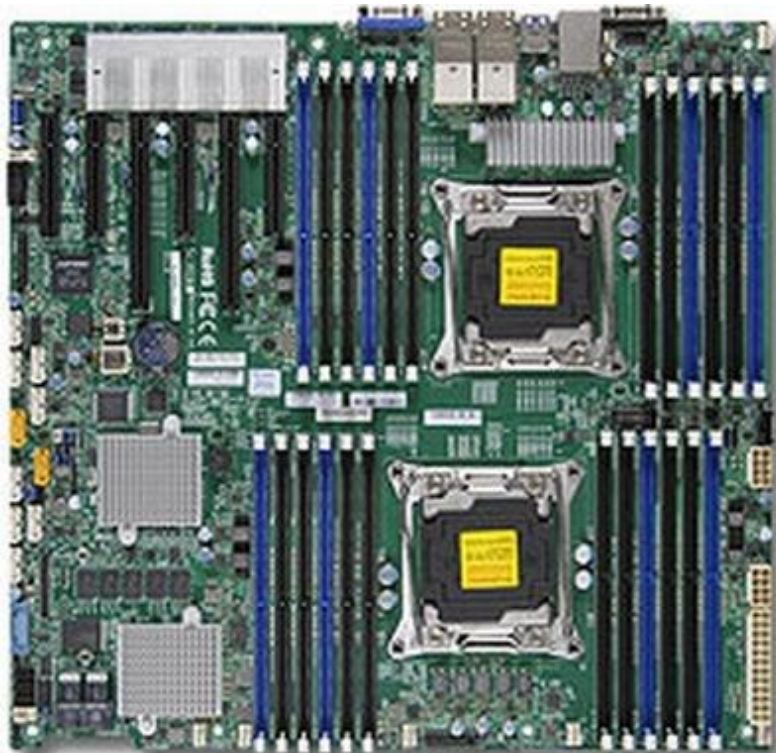
---

- In computing, essentially all (time/energy) cost is due to moving data over a distance.
  - *Moving data into the CPU*
  - *Moving data through the CPU*
- Fundamental limits to shrinking distances:
  - *Quantum effects soon to take over (one cannot work with less than an electron; also noise)*
  - *Failure of Dennard scaling.*
  - *We have only 2/3 dimensions to pack stuff into space. (3d chips, but also consider cooling!)*
- Communication links need space: Buses, networks are bottlenecks
- In many (most?) applications, CPUs are mainly waiting for data to arrive.
  - *Not just explicit data management applications.*
  - *Cache hierarchies, prefetching, ...*



# Nodes, sockets, CPUs, RAM, I/O etc.

---

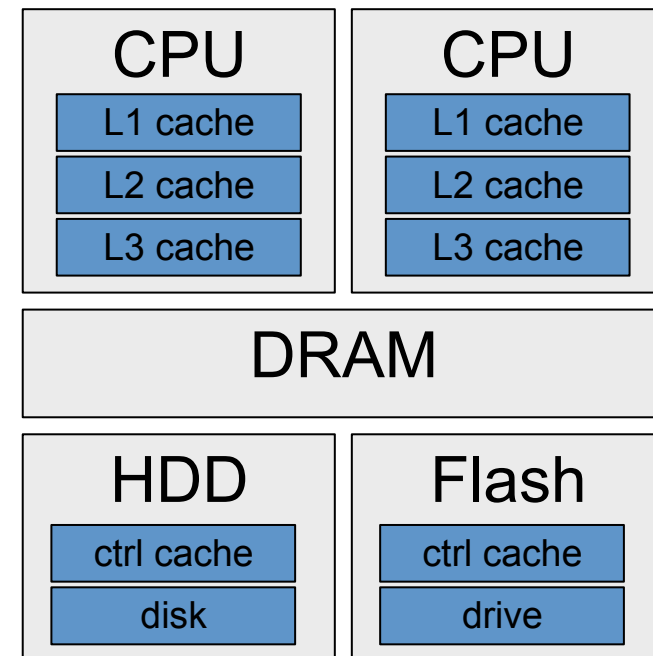


# Memory hierarchies; time scales

- Core i7 Xeon 5500 Series Data Access Latency (approximate)

- L1 CACHE hit, ~4 cycles
- L2 CACHE hit, ~10 cycles
- L3 CACHE hit, line unshared ~40 cycles
- L3 CACHE hit, shared line in another core ~65 cycles
- L3 CACHE hit, modified in another core ~75 cycles
- remote L3 CACHE ~100-300 cycles (@~ 0.3ns/cycle)

- Local DRAM ~60 ns
- Remote DRAM ~100 ns
- Accessing a hard drive – 10ms(seek)
  - 0.1ms/page (transfer)
- Accessing a tape – minutes (seek)
- L1 to DRAM:  $10^2$ x; DRAM to HD:  $10^5$ x slowdown



# Caches

---

- The further away a cache is from the core, the larger, slower, and cheaper (due to different technologies used) it is.
- While Moore's law was active, the speed of computation (and register access) was diverging from the speed of RAM – more and more cache layers were inserted.
  - *When does it make sense to introduce another layer of cache?*

# Units of computation

---

- Nodes, sockets (CPUs), and cores. Why so many abstractions?
- Data locality considerations!
- Peers of these three types communicate using very different technologies, which have vastly different latencies. The argument for having all of these is the same as for memory hierarchies!

# Data center servers vs. supercomputers

---

- Supercomputers have higher compute/network density, not just by volume/watt but also per \$. Network fabrics are quite different from those in data centers; there is much higher connectedness.
- Data center servers support better data locality for large data volumes; they have more RAM per node and usually have their own secondary storage devices.
- Use case: Matrix computations.
  - *Supercomputers better for dense linear algebra*
  - *Data center servers often more cost-effective for certain sparse linear algebra computations and tasks (such as data cleaning and transformation) where the amount of data moved is high relative to the intensity of computation.*

# Locality Principles

---

- Caching
  - *Keep a working set of data that is used frequently close.*
- Prefer sequential over random access
  - *Physics governs that some data is better read/written sequentially than by random access.*
- Partitioning
  - *Some tasks allow to consider data in parts: either use all resources to process a part at a time, or work on the partitions in parallel.*
- Use cases: Out-of-core algorithms: joins, sorting (of data on hard drives)



# Caching

---

- Ubiquitous in systems
  - *CPU caches*
  - *MMUs: TLB*
  - *Networks (edge caches)*
  - *OS/DBMS buffers; storage device controller caches (hard drives, solid state, ...)*
  - *DBMS: materialized views*
- Caching is frequently (erroneously) equated with locality. [Salzer and Kaashoek]
- Prefetching: Speculative filling of cache – usually assumes sequential access
  - *CPU branch prediction, storage device controllers – what else could be done?*

# Sequential access

---

- Sequential access faster than random access
  - *Hard drives*
    - *Mechanically moving parts: seek time  $\gg$  transfer time*
    - *Reading a byte is not cheaper than reading a page*
  - *Flash/solid state: only large blocks can be written, only very large ones erased.*
  - *DRAM*
    - *Block addressing and transfer via the bus*
    - *TLBs (again)*
- Examples: Block nested loops join, external-memory sort





# Partitioning

---

- Decomposability and embarrassing parallelism.
  - *When can different parts of the data be processed completely independently? (No communication needed)*
  - *Map/reduce*
- But: not applicable everywhere.
  - *Frequent: the graph of our data (dependencies, relationships) has small diameter even though sparse (“small world phenomenon”)*
  - *Such graphs have no small cuts (see next video)*
- Out-of-core algorithm examples: External memory sort, GRACE hash join



# Locality and data structure(s)

---

Christoph Koch

*School of Computer & Communication Sciences, EPFL*

# Data Structures: Arrays

---

- Does the storage layout match the looping order of the algorithms that access the array? – sequential access vs random access.
- Example: Matrix
- Stored as A11, A12, ..., A1n, A21, A22, A2n, ..., Amn
- Loop: for i in 1 ... n { for j in 1 ... m { Aij ... }} efficient
- Loop: for j in 1 ... m { for i in 1 ... n { Aij ... }} inefficient ( due to block addressing, prefetching, cache misses)
- Align storage layout with use cases (pattern of access) if possible, or vice versa.
  - *Loop reordering in compilers.*
  - *Sorting, nesting, co-clustering in DBMS.*

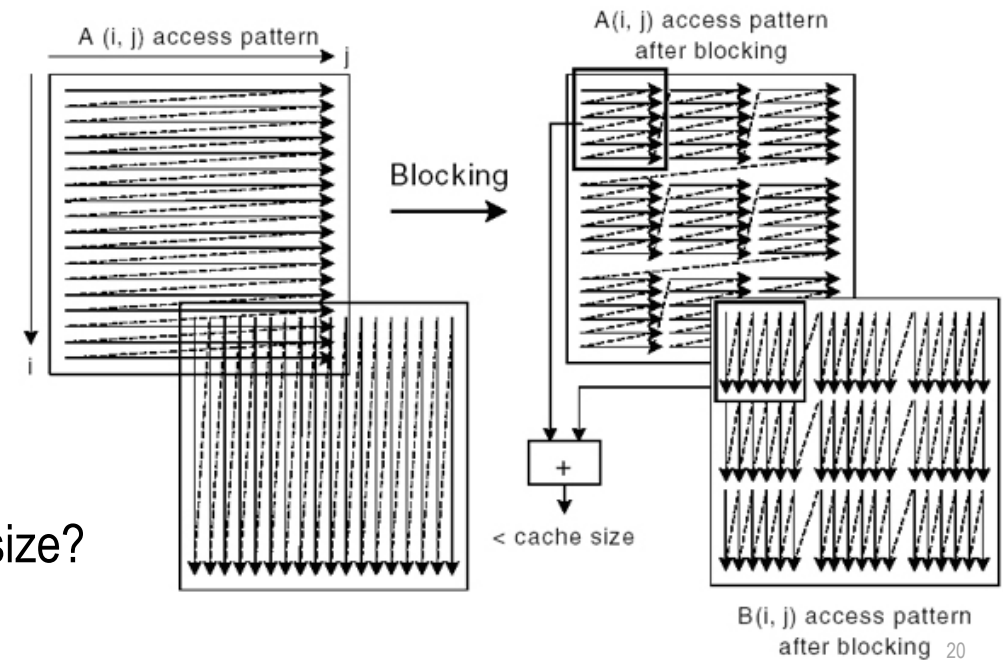
A11	A21	...	Am1
A12	A22	...	Am2
...	...		...
A1n	A2n	...	Amn

# Loop tiling

- Assume (2d-)array to be processed is larger than cache:
- Reorder loops to process large array by small array regions (tiles).

```
for (i=0; i< MAX; i++) {  
    for (j=0; j< MAX; j++) {  
        A[i,j] = A[i,j] + B[j, i];  
    }  
}
```

New loop structure? Tile size vs. cache size?  
(source: intel)



# Arrays/Relations: Row vs. columnar representation

- (OLAP) databases: Row vs column-stores
  - Many queries use only some columns of a relational table. Only fetch the data from disk that you need.
  - Better on-disk compression of columns
    - E.g. many consecutively stored phone numbers compress better than random data.
  - Lots of hype about this!
- OO PLs, e.g. Java (VM)
  - row representation: array of struct  $\{ \text{int } a, \text{int } b \}$
  - column representation: struct  $\{ \text{int } a[], \text{int } b[] \}$
  - column representation much more efficient:
  - Much fewer objects created ( $O(1)$  vs  $O(\text{array size})$ ).  
Boxing/unboxing overhead!



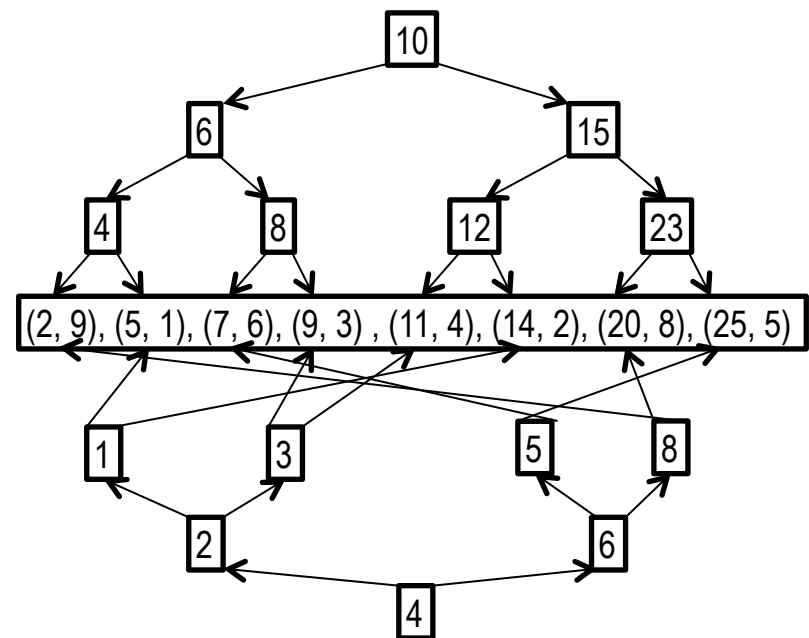
# Data Structures: Trees

---

- Example: Balanced binary tree.
- On each level, twice as many nodes as in the level above.
  - *Exponential growth.*
- There is NO WAY to store data in linear (or finite-dimensional) memory so that parent-child pairs remain close!
- But: can keep siblings local: breadth-first enumeration.
- Or: leaf level size dominates: leaf level of depth-first enumeration is essentially local.
  - *Basis of indexing!*

# Tree indexes

- The leaf level of a balanced tree is (essentially) local in reasonable representations.
- Idea of primary (B-)tree indexes in DBMS: leaf level is aligned with sort order in data file.
- Range lookup.
  - *Find first matching element using index, then scan data file sequential until range condition becomes false.*
  - *But: no two different tree indexes can index locally into shared data! (Secondary indexes)*
  - *Thus indexes are not as effective as often naively assumed.*
- Even prim. index is not very useful: same log-many random accesses as in binary search on data file.



# Types of graphs

---

- Graphs with edge relations that have a local representation
  - *Simple, special deterministic constructions: chains, trees – not interesting here.*
- Graphs with (relatively) small cuts (partitioning/parallelization!)
  - *Resource-constrained graphs:*
    - *(Almost) planar embeddings into a 2d surface (map)*
    - *Constraints on density of edges crossing any line on the map.*
    - *Not trees, but relatively low degree of cyclicity. Few (redundant) nonlocal links*
  - *Road networks – roads take space, too many is pointless.*
  - *Physical internetworks – deal with line cost, routing complexity.*
  - *The brain (in theory)*
- All other graphs: locality nightmare.
  - *Internet communication patterns, social networks, the brain (in practice), ...*



# Types of graphs

---

- Random graphs: edges added randomly
  - *Above edge-to-node ratio 1, essentially all of the graph is connected (monster connected component)*
  - *Have low degrees of separation (small world phenomenon) and NO SMALL CUTS already when sparse (linear edge-to-node ratio)*
- Real-world graphs and networks
  - *Exclude resource-constrained graphs, e.g. road networks, physical internetworks*
  - *Essentially all other graphs/networks: internet communication patterns, social networks, the brain, ...*

# Real-world graphs

---

These notions are essentially equivalent/interchangeable:

- Power law graphs: random graphs in which #neighbors follows a power law.
- Social networks: “rich get richer” phenomenon: popular nodes are more likely to get further connections.
- Small world graphs: “k degrees of separation”
  - *Example construction (Kleinberg): Take grid (fishing net) and add a certain ratio of non-local shortcuts.*
- Already for sparse graphs
- Deterministic construction: expander graphs

# Data Structures: Graphs and Networks

---

- (Non-resource bounded) real graphs have no small cuts:
  - *Take a graph. There is no partition of the nodes into two about equally sized sets such that the #edges crossing the partition is less than linear.*
  - *CANNOT be partitioned effectively to handle regions independently without requiring lots of “communication” between regions.*
  - *Essentially impossible to parallelize graph analytics effectively.*
  - *Everyone still does it. Horrible performance (Pregel, Giraffe) – every node has to talk to every node in every step.*
    - *A worst-case scenario from the locality perspective.*
- But: small-world phenomenon
  - *There are short paths between any two nodes (routing!)*
  - *Does not mean communication is spatially local, but is local if you only count hops!*
  - *Theory of weak ties (sociology) vs. routing heuristics!*

## Graphs with small cuts    [not relevant to the final exam]

---

Definition: A graph  $G=(V,E)$  has tree-width  $(\leq)k$  if there is a set  $H$  (the hypergraph/tree decomposition of  $G$ ) of subsets of  $V$  such that,

- for each  $S$  in  $H$ ,  $|S| \leq k+1$  and
  - for each edge  $\{v_1, v_2\}$  in  $E$ , there exists at least one  $S$  in  $H$  such that  $v_1, v_2$  in  $S$  ( $H$  covers  $E$ ).
- 
- Test if graph has tree-width  $k$ : Bodlaender's algorithm (linear in  $|E|$  if  $k$  is fixed)
  - Compute tree-width of graph: NP-complete.
  - NP-complete problems: fix tree-width of combinatorial structure  $\Rightarrow$  in P.
  - Traverse tree decomposition of graph is often the best-known technique for processing hard combinatorial problems.

## Classroom exercise: Distributed All-Pairs Shortest Paths

---

- Floyd-Warshall:  
     $sp(i, j, 0) := w(i, j)$   
     $sp(i, j, k) := \min(sp(i, j, k-1), sp(i, k, k-1) + sp(k, j, k-1))$
- On a road network



# Out-of-core algorithms: external sorting

---

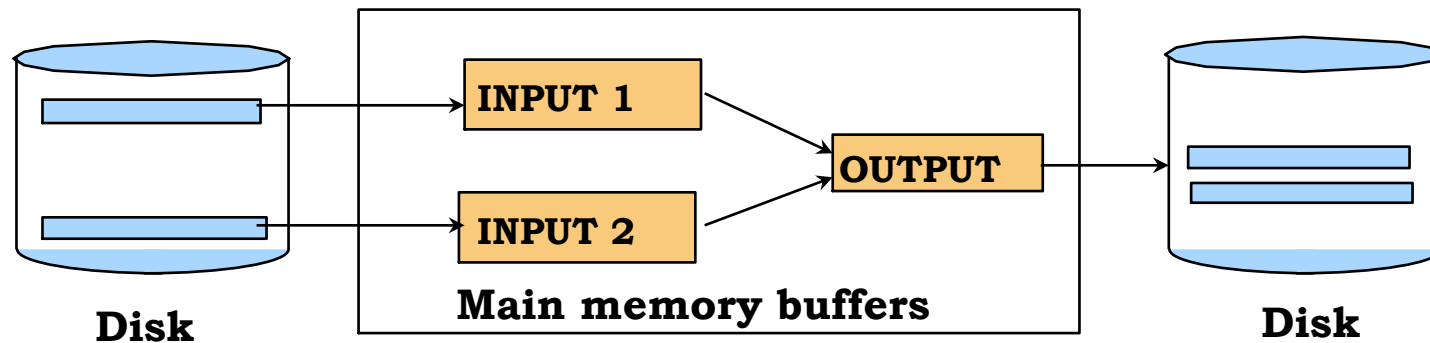
Christoph Koch

*School of Computer & Communication Sciences, EPFL*

## 2-Way Sort: Requires 3 Buffers

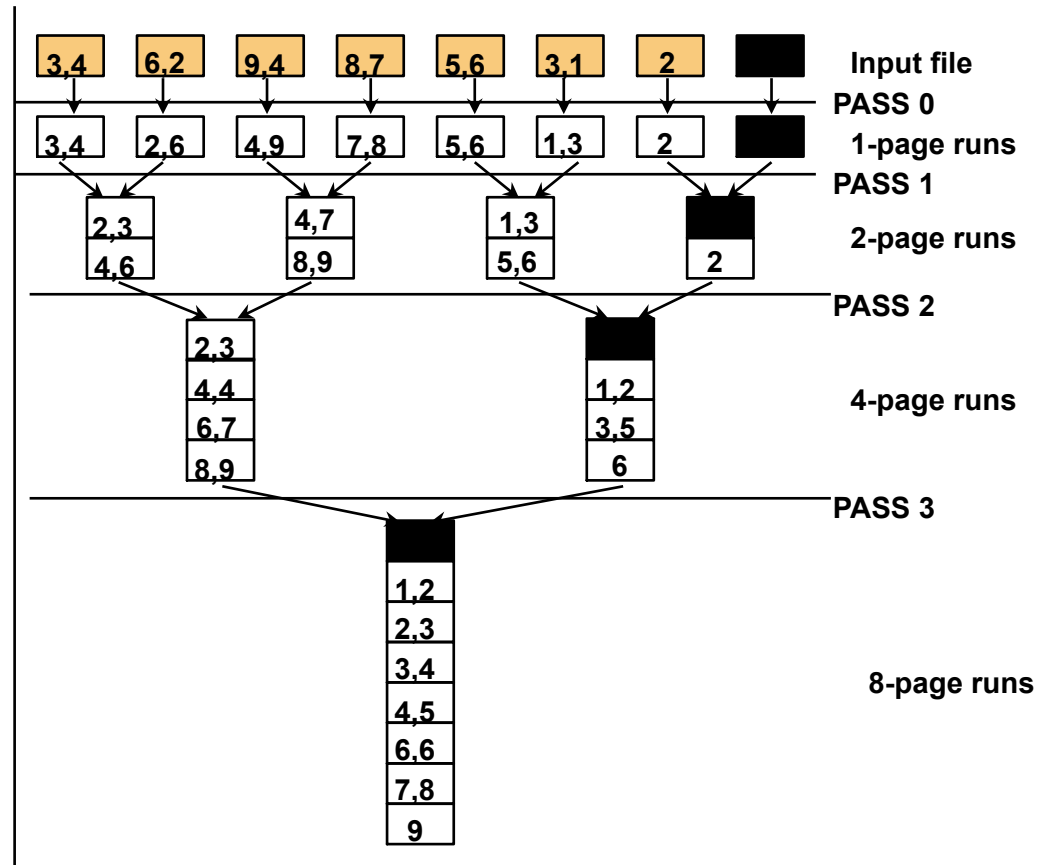
---

- Pass 1: Read a page, sort it, write it.
  - *only one buffer page is used*
- Pass 2, 3, ..., etc.:
  - *three buffer pages used.*



# Two-Way External Merge Sort

- In each pass we read + write each page in file.
- $N$  pages in the file  $\Rightarrow$  the number of passes  $= \lceil \log_2 N \rceil + 1$
- So the total cost is  $2N(\lceil \log_2 N \rceil + 1)$
- Idea: **Divide and conquer:** sort subfiles and merge

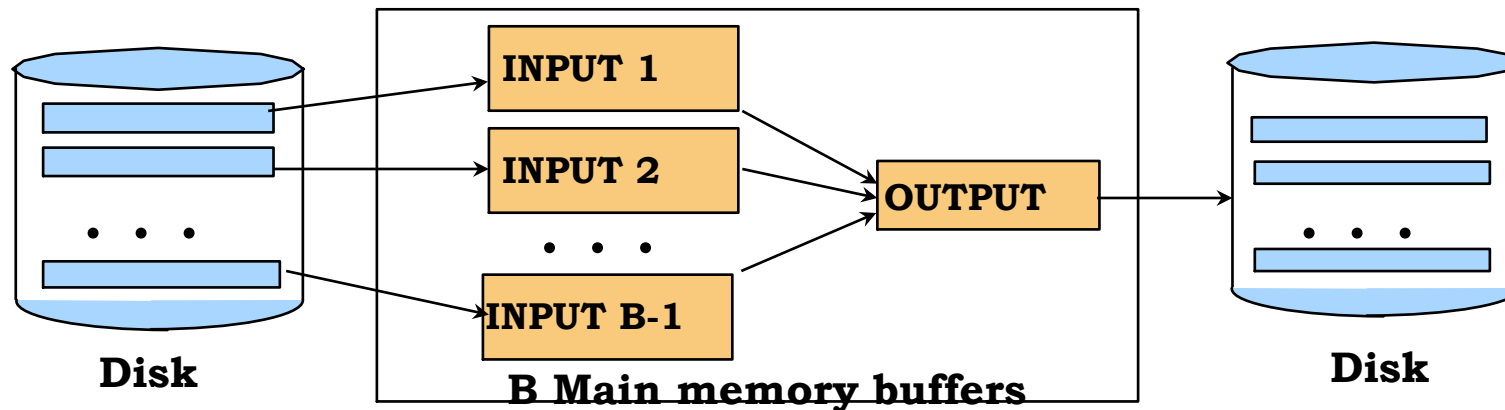




# General External Merge Sort

---

- To sort a file with  $N$  pages using  $B$  buffer pages:
  - *Pass 0: use  $B$  buffer pages. Produce  $\lceil N / B \rceil$  sorted runs of  $B$  pages each.*
  - *Pass 2, ..., etc.: merge  $B-1$  runs.*



## Cost of External Merge Sort

---

- Number of passes:  $1 + \lceil \log_{B-1} \lceil N / B \rceil \rceil$
- Cost =  $2N * (\text{\# of passes})$
- E.g., with 5 buffer pages, to sort 108 page file:
  - Pass 0:  $\lceil 108 / 5 \rceil = 22$  sorted runs of 5 pages each (last run is only 3 pages)
  - Pass 1:  $\lceil 22 / 4 \rceil = 6$  sorted runs of 20 pages each (last run is only 8 pages)
  - Pass 2: 2 sorted runs, 80 pages and 28 pages
  - Pass 3: Sorted file of 108 pages

## Number of Passes of External Sort

---

N	B=3	B=5	B=9	B=17	B=129	B=257
100	7	4	3	2	1	1
1,000	10	5	4	3	2	2
10,000	13	7	5	4	2	2
100,000	17	9	6	5	3	3
1,000,000	20	10	7	5	3	3
10,000,000	23	12	8	6	4	3
100,000,000	26	14	9	7	4	4
1,000,000,000	30	15	10	8	5	4

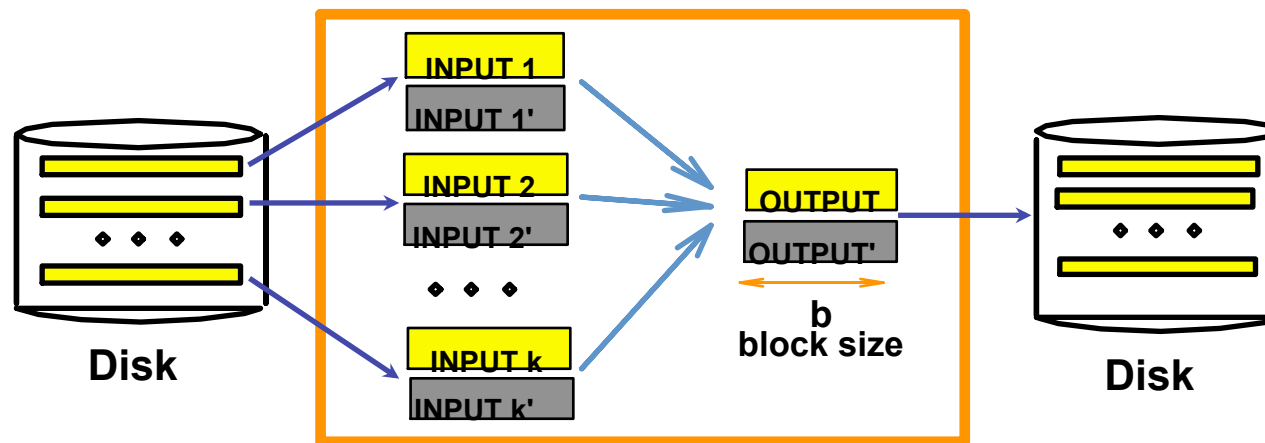
## I/O for External Merge Sort

---

- ... longer runs often means fewer passes!
- Actually, do I/O a page at a time
- In fact, read a *block* of pages sequentially!
- Suggests we should make each buffer (input/output) be a *block* of pages.
  - *But this will reduce fan-out during merge passes!*
  - *In practice, most files still sorted in 2-3 passes.*

# Double Buffering

- To reduce wait time for I/O request to complete, can *prefetch* into *'shadow block'*.
- Potentially, more passes; in practice, most files *still* sorted in 2-3 passes.



B main memory buffers, k-way merge



# Out-of-core algorithms: joins

---

Christoph Koch

*School of Computer & Communication Sciences, EPFL*

## Simple Nested Loops Join

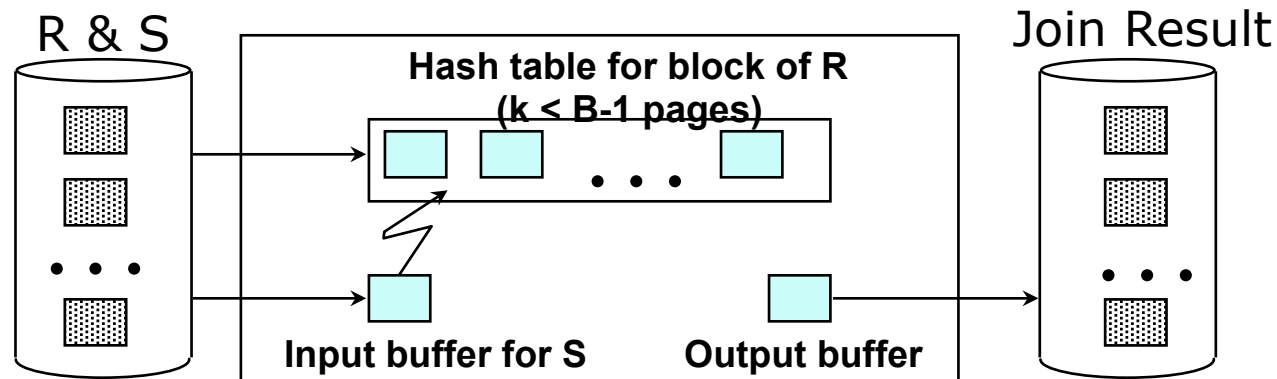
---

```
foreach tuple r in R do
    foreach tuple s in S do
        if  $r_i == s_j$  then add  $\langle r, s \rangle$  to result
```

- For each tuple in the *outer* relation R, we scan the entire *inner* relation S.
  - *Cost:*  $M + p_R * M * N = 1000 + 100 * 1000 * 500$  I/Os.
- Page-oriented Nested Loops join: For each *page* of R, get each *page* of S, and write out matching pairs of tuples  $\langle r, s \rangle$ , where r is in R-page and S is in S-page.
  - *Cost:*  $M + M * N = 1000 + 1000 * 500$
  - If smaller relation (S) is outer, cost =  $500 + 500 * 1000$

## Block Nested Loops Join

- Use one page as an input buffer for scanning the inner S, one page as the output buffer, and use all remaining pages to hold ``block'' of outer R.
- *For each matching tuple  $r$  in R-block,  $s$  in S-page, add  $\langle r, s \rangle$  to result. Then read next R-block, scan S, etc.*





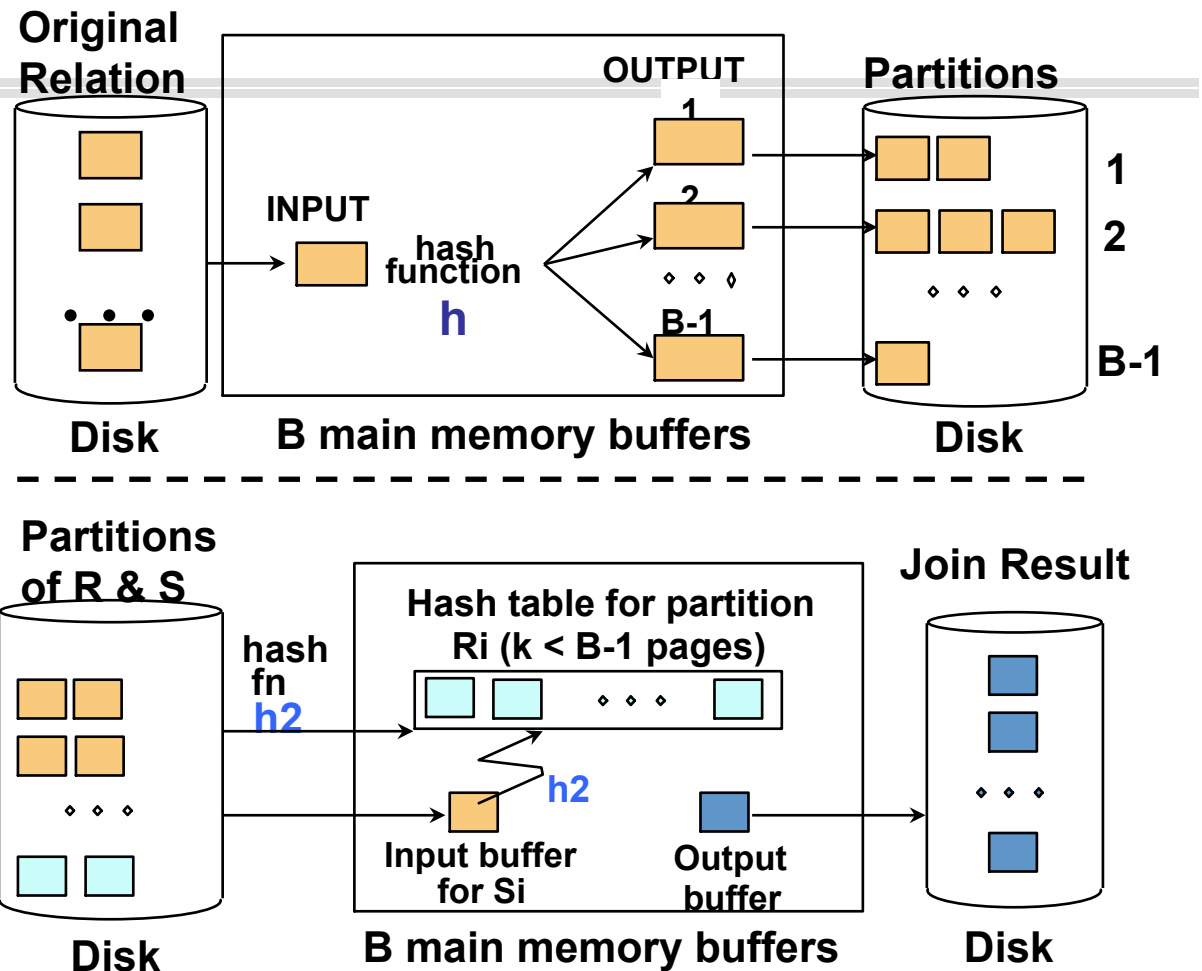
## Block Nested Loops Join

---

- Cost: Scan of outer + #outer blocks \* scan of inner
  - $\#outer\ blocks = \lceil \# of\ pages\ of\ outer / blocksize \rceil$
- With sequential reads considered, analysis changes: may be best to divide buffers evenly between R and S.
  - *Depends on whether join processing can keep up with the scan of the inner relation.*

# Hash-Join

- Partition both relations using hash fn **h**: R tuples in partition i will only match S tuples in partition i.
- Read in a partition of R, hash it using **h2** ( $\neq h$ ). Scan matching partition of S, search for matches.



## Observations on Hash-Join

---

- #partitions  $k < B-1$ , and  $B-2 > \text{size of largest partition}$  to be held in memory. Assuming uniformly sized partitions, and maximizing  $k$ , we get:
  - $k = B-1$ , and  $M/(B-1) < B-2$ , i.e.,  $B$  must be  $> \sqrt{M}$
- If we build an in-memory hash table to speed up the matching of tuples, a little more memory is needed.
- If the hash function does not partition uniformly, one or more  $R$  partitions may not fit in memory. Can apply hash-join technique recursively to do the join of this  $R$ -partition with corresponding  $S$ -partition.

## Cost of Hash-Join

---

- In partitioning phase, read+write both relns;  $2(M+N)$ .
- In matching phase, read both relns;  $M+N$  I/Os.
- Sort-Merge Join vs. Hash Join:
  - *Given a minimum amount of memory, both have a cost of  $3(M+N)$  I/Os.*
  - *Hash Join superior on this count if relation sizes differ greatly.*
  - *Also, Hash Join shown to be highly parallelizable.*
  - *Sort-Merge less sensitive to data skew; result is sorted.*